
Optimization of Maintenance Operations for Urban EV Charging Infrastructure under Uncertainty

Philipp Hümmer

2485927



Master's Thesis

Chair of Business Informatics and Business Analytics
University of Würzburg

Advisor/Supervisor: Tim Lachner
Chair: Prof. Dr. Christoph Flath

Würzburg, den 23.07.2026

Contents

List of Figures	III
List of Tables	IV
1 Introduction	1
2 Problem Definition	3
2.1 Operational Setting	3
2.2 Environment and Disturbances	4
2.3 Assumptions and Cost Model	5
2.4 Motivating Example	6
3 Literature Review	9
3.1 Maintenance Optimization in Infrastructure Systems	9
3.1.1 Concept and Objectives of Maintenance Optimization	9
3.1.2 Maintenance Strategies: Preventive, Reactive, Condition-Based	10
3.1.3 Uncertainty, Availability, and Cost as Central Decision Variables	10
3.1.4 Transition to Mobile Maintenance and Service Interventions	11
3.1.5 EV Charging Infrastructure as a Maintenance Context	12
3.2 Stochastic Vehicle Routing Problems and Dynamic Service Planning	12
3.2.1 Basic Concepts of VRP and DVRP	13
3.2.2 Uncertainties in Routing: Requests, Service Times, Travel Times, Capacities	13
3.2.3 Modeling Approaches: MDP, Multi-Stage Stochastic Programming, Look-ahead	14
3.2.4 Solution Families: Exact Methods, Heuristics, Metaheuristics	14
3.2.5 Route-based MDPs and Their Relevance for This Thesis	15
3.2.6 Limits of Classical SDVRP Models for Maintenance Tasks	16
3.2.7 Spatial Clustering as a Pre-Structuring Mechanism	16
3.3 Approximate Dynamic Programming for Dynamic Routing Problems	17
3.3.1 Motivation: Why Exact Dynamic Programming Does Not Scale	17
3.3.2 Basic Idea of ADP: States, Actions, Value Functions	17
3.3.3 Important ADP Approaches: Rollout, VFA, and Offline-Online ADP	18
3.3.4 Distinction from Reinforcement Learning	19
3.3.5 Relevance of ADP for Dynamic Maintenance Decisions	19
4 Theoretical Framework	20
4.1 Markov Decision Process	20
4.1.1 Sets, indices, and temporal structure	20
4.1.2 State representation	21
4.1.3 Action space	22

4.1.4	Transition dynamics	23
4.1.5	Reward function	24
4.2	State-Action-Value Structure	25
4.3	ADP Foundations and the Offline-Online Approach	26
5	Solution Methodology	28
5.1	Shared Zone Selection Logic	28
5.2	Myopic	30
5.3	Myopic+	31
5.4	Cost Function Approximation	32
5.5	Dynamic Balance	33
5.6	Value Function Approximation	35
6	Experimental Setup	36
6.1	Simulation Design	37
6.2	Routing and Travel Times	38
6.3	Data Basis	38
7	Results and Analysis	40
7.1	Policy Mechanism Comparison	40
7.2	Impact of Zone Selection on Cost	41
7.3	VFA Rollout: Performance and Limitations	46
7.4	Statistical Validation	50
8	Discussion	53
8.1	Interpretation of Results	53
8.2	Limitations	54
8.3	Practical Relevance and Transferability	55
9	Conclusion and Outlook	57
9.1	Summary of Key Findings	57
9.2	Contributions of This Thesis	58
9.3	Open Research Questions and Future Directions	58
9.4	Concluding Remarks	59
	Bibliography	62
	List of Acronyms	63
A	Appendix A	64

List of Figures

1	Spatial distribution of the 397 charging stations across the Würzburg urban service area, with the central depot marked.	4
2	Two initial plan strategies for the same set of stations. Left: routing-efficiency approach visits A second despite low urgency. Right: maintenance-urgency approach defers A to last, ensuring critical stations are covered first.	7
3	Two replanning decisions under the same incoming disruption E. Left: dropping B (largest detour) leaves the most critical station unserved. Right: dropping A (recently serviced, low power) allows B to be serviced and averts failure risk.	8
4	Hierarchy of decision policies developed in this thesis. All four base policies share the same zone-selection layer (Section 5.1) and differ only in their routing score and drop score, shown in each box. The VFA layer (Section 5.6) is not a fifth base policy but an optional rollout wrapper that can be placed on top of any of the four.	29
5	Execution timeline of a single simulated day for both teams (CFA, value-based zone selection, seed 1, day 1), reconstructed from the actual simulation log. Both teams start their routine route at 8:00; disruptions (Type 1 and Type 2) are inserted mid-route via cheapest insertion, and the shared 60-minute lunch break is fixed at 12:00–13:00 and interrupts the active service stop if one is ongoing. Of the 20 routine tasks planned for the day, 10 could not be completed and were carried over to day 2.	37
6	Three-phase benchmark for the number of zones k : mean CFA total cost by zone count. Each point represents the average over 100 simulation runs; colours indicate the benchmark phase.	39
7	Mean service day per zone for CFA under centrality-based (left) and value-based (right) zone selection, averaged over 500 runs. Green = serviced early, red = serviced late. Under centrality, the service order is strictly concentric, with centrally located zones always visited first and peripheral zones last. Value-based selection breaks this pattern, pulling economically critical zones forward to early service days regardless of their position in the network.	43
8	Total cost distributions for all four policies under centrality (light) and value-based (dark) zone selection (500 runs each). Boxes show interquartile range, and whiskers extend to $1.5 \times \text{IQR}$	44
9	Cumulative costs (top) and daily disruption counts (bottom) for Myopic and CFA under value-based zone selection, seed 497 (single run). The dashed vertical line marks day 9, where the gap widens most sharply. Myopic total: 52,491 €; CFA total: 37,191 €.	45

10	Distribution of paired cost differences (base minus rollout) per run for each policy, shown as histograms with 250€ bins. Positive values indicate the rollout lowered cost, and negative values indicate the rollout increased cost. The dashed vertical line marks zero. Runs with $\Delta = 0$ (no override was triggered) are excluded from the histogram and reported separately in the annotation. The x-axis is clipped to $\pm 4,000$ €, and a small number of extreme values fall outside the plotted range. All variants use value-based zone selection (500 runs each).	47
11	Distribution of replan overrides per run for all four rollout variants (500 runs each). Dashed vertical line shows the mean.	48
12	Cohen's d for all key comparisons. Dashed vertical lines mark conventional thresholds (small: 0.20, medium: 0.50, large: 0.80). Zone-selection effects are the only comparisons with a large effect size, and all other comparisons fall below the medium threshold ($d = 0.50$).	52

List of Tables

1	Policy comparison under centrality-based zone selection (500 runs each). SD refers to the standard deviation of total costs across runs.	41
2	Cost comparison across zone-selection modes for all four policies (500 runs each). Δ is computed relative to centrality for each policy.	42
3	Rollout effect per policy. All variants use value-based zone selection (500 runs each). Override rate = mean replan overrides / mean total carryover (≈ 48).	46
4	Seed 6 cascade: cumulative total costs around the single day-11 override (Myopic with value-based zone selection, base vs. rollout run).	49
5	Paired statistical tests for key comparisons (500 matched seeds each). Cohen's d is the primary criterion for practical relevance, and p -values are from the paired t -test and Wilcoxon signed-rank test.	51
6	Summary of notation used in Chapter 4	64

Abstract

Public electric vehicle charging networks depend on the reliability of their stations, yet maintaining hundreds of geographically distributed chargers is complicated by uncertainty, since faults occur stochastically, vary in severity, and compete with regular preventive inspections for the limited daily capacity of a small number of maintenance teams. This thesis studies how such maintenance operations can be planned dynamically under uncertainty, using the charging network of Würzburg, with 397 stations serviced by two teams, as a case study.

The problem is formulated as a route-based Markov decision process in which routing decisions, rather than individual node selections, constitute the unit of action, and in which both an initial daily plan and subsequent intraday replanning must be represented. Since an exact solution is computationally intractable at this scale, the thesis develops a hierarchy of approximate dynamic programming policies of increasing sophistication. Myopic is a purely distance-driven greedy baseline, and Myopic+ weights routing decisions by station power. Cost Function Approximation (CFA) learns a station-specific value function via contrastive suffix-simulation regression and uses it to score routing and drop decisions, while Dynamic Balance extends CFA with a state-dependent parameter trading off economic value against routing efficiency. A Value Function Approximation (VFA) layer augments any of these policies with a Monte Carlo rollout at critical drop decisions. All policies share a common K-Means-based zone-selection mechanism, which is evaluated both in a geography-only and a value-based variant.

The policies are compared using a discrete-event simulation calibrated to real station data and time-dependent travel times from OSRM and Google Distance Matrix data, over 500 Monte Carlo runs per configuration with seed-matched disturbance sequences. The results show a clear hierarchy of effects. The choice of zone-selection strategy is the dominant cost driver, reducing total cost by 14–16 % ($d = 1.29\text{--}1.50$), far outweighing differences between routing policies ($d \leq 0.37$) and the benefit of the rollout-based VFA layer ($d \leq 0.08$, not economically meaningful). Under a fixed, non-value-based zone selection, all four policies perform statistically indistinguishably, whereas once zone selection is value-based, the simplest policy (Myopic) falls measurably behind the value-weighted policies, while CFA and Dynamic Balance remain effectively equivalent to each other. The practical implication is that a well-chosen zone assignment with a simple routing policy outperforms a poorly chosen zone assignment paired with the most sophisticated policy developed in this thesis, suggesting that investment in zone-selection quality yields substantially higher returns than further refinement of the routing heuristic itself.

1 Introduction

The transition toward electromobility has led to a rapid expansion of public charging infrastructure in cities across Germany, and Würzburg is no exception. As the number of publicly accessible charging stations grows, so does the operational burden of keeping this infrastructure functional. Every charging station that is out of service represents not merely a technical fault but a direct loss of capacity for a network whose entire value proposition rests on availability, since a driver who arrives at a broken charger loses trust in the network as a whole, and an operator who lets faults linger risks both revenue and reputation. Maintaining hundreds of geographically distributed stations reliably is therefore not a peripheral concern but a core operational challenge for any charging network operator.

In practice, this challenge is compounded by uncertainty. Two maintenance teams cannot know in advance which stations will fail on a given day, when a failure will be reported, or how severe it will be. Some faults are minor and quickly resolved, while others require a station to be dismantled, transported, and reassembled. Regular preventive inspections must be scheduled alongside this unpredictable corrective workload, and both compete for the same limited pool of team-hours, vehicles, and working time. A route plan that looks efficient at 8:00 may be obsolete by 10:00, once a new disturbance forces a team to deviate from its planned path. Any planning approach that ignores this dynamic, sequential character of the problem, treating it as a single static routing exercise solved once at the start of the day, is bound to misrepresent how maintenance actually unfolds in practice.

Classical vehicle routing models are well suited to static, fully known problem instances but are conceptually mismatched to this setting, since they neither represent the sequential resolution of uncertainty nor the fact that today's routing decisions influence which stations remain at risk tomorrow. Markov Decision Processes offer a theoretically rigorous alternative that explicitly models this sequential, uncertain structure, but an exact MDP formulation of a problem with 397 stations, two teams, and stochastic disturbances arriving throughout the day is computationally intractable, since the state and action spaces are simply too large to enumerate or solve exactly. This tension between the conceptual adequacy of the MDP framework and its practical intractability motivates the central question addressed in this thesis:

How can maintenance operations for EV charging stations be planned dynamically under uncertainty, so as to balance downtime, travel costs, and resource utilization efficiently?

This thesis addresses this question by formulating the maintenance problem as a route-based MDP and then solving it approximately through a hierarchy of policies grounded in approximate dynamic programming (ADP). Rather than attempting to solve the MDP exactly, the thesis develops a family of increasingly sophisticated greedy heuristics, ranging from a purely distance-driven baseline to policies that learn an explicit economic value for each station and combine it with a rollout-based lookahead, and evaluates them against

each other in a realistic, data-driven simulation of the Würzburg charging network. Guided by this overall aim, the thesis makes four main contributions. First, it reviews the relevant literature on maintenance optimization, stochastic and dynamic vehicle routing, and ADP, situating the problem at the intersection of these three research streams. Second, it formulates the operational maintenance problem as a route-based MDP, with an explicit state representation, action space, transition dynamics, and cost structure tailored to the charging-infrastructure context. Third, it develops and implements a hierarchy of ADP-based decision policies, namely Myopic, Myopic+, Cost Function Approximation (CFA), Dynamic Balance (DB), and a rollout-based Value Function Approximation (VFA) layer, that translate this formulation into operational routing and replanning rules.¹ Fourth, it empirically evaluates these policies on real station data for Würzburg’s 397 charging stations, enriched with a simulated stochastic failure process, across 500 Monte Carlo runs per configuration, isolating the effects of routing policy, zone-selection strategy, and rollout-based lookahead on total operating cost.

The remainder of this thesis is organized as follows. Chapter 2 introduces the operational planning problem in detail, describing the service network, the nature of disturbances, the cost model, and a motivating example that illustrates why state-aware prioritization matters. Chapter 3 reviews the relevant literature on maintenance optimization, stochastic and dynamic vehicle routing, and ADP. Chapter 4 formalizes the problem as a route-based MDP and develops the state-action-value perspective that underlies the ADP approach. Chapter 5 presents the solution methodology, describing the shared zone-selection mechanism and the hierarchy of policies from Myopic to the VFA rollout layer. Chapter 6 details the experimental setup, including the simulation design, travel-time model, and data basis. Chapter 7 presents the results, comparing policy mechanisms, examining the impact of zone selection, evaluating the VFA rollout, and providing statistical validation of the key findings. Chapter 8 discusses the interpretation, limitations, and practical relevance of these results, and Chapter 9 concludes with an outlook on future work.

¹The full source code is publicly available at <https://github.com/philipphuemmer/masterarbeit>.

2 Problem Definition

This chapter introduces the practical planning problem addressed in the thesis, namely the maintenance and disturbance handling of a network of electric vehicle charging stations.

The problem involves the daily routing and scheduling of mobile maintenance teams responsible for servicing a geographically distributed network of charging stations. The planning problem is not static, since in addition to regular preventive tasks, new disturbance events may occur during the day and require operational responses on short notice. As a result, the system combines elements of preventive maintenance, corrective maintenance, and dynamic routing.

The decision process is organized on a daily basis. At the beginning of each day, an initial service plan is constructed for both teams. During the operating day, this plan may be revised in response to newly occurring faults.

2.1 Operational Setting

The operational scenario considered in this thesis is based on the charging infrastructure of the city of Würzburg, consisting of 397 stations distributed across the urban service area. These charging stations form the service network that must be maintained by two mobile service teams. Each team starts its route at a central depot and is required to return to the depot by the end of the working day. The service region is represented as a graph whose node set consists of the depot and all charging stations. Travel between stations is therefore a core element of the problem, since every maintenance action requires physical movement through the network.

The planning horizon is structured as a daily operational process. The working day runs from 8:00 to 17:00. At the beginning of the day, a full initial plan is generated. During the day, replanning decisions are triggered event-driven on an hourly grid whenever new disturbance events arrive. This means that routing decisions are not made once and then executed unchanged. Instead, they are revised as new information becomes available. Such a planning logic is essential in practice because maintenance operations in charging infrastructure are rarely fully predictable at the start of the day.

A further operational characteristic is the presence of a deterministic lunch break between 12:00 and 13:00. In the route-based setup, this break is not treated as an independent decision variable. Instead, it is inserted according to a fixed operational rule and shifts all subsequent arrival and departure times accordingly. This reflects the idea that some daily routines are known in advance and need not be optimized explicitly, but still influence route feasibility and time consumption.

The spatial structure of the system also plays an important role. Because the charging stations are distributed across the service region, route efficiency depends strongly on travel distances, travel times, and geographic clustering. For this reason, the operational logic uses zones as a practical pre-structuring device. The stations are grouped into zones which




Figure 1: Spatial distribution of the 397 charging stations across the Würzburg urban service area, with the central depot marked.

support the initial selection of promising service areas for the two teams. These zones reduce routing complexity and guide the daily assignment of work areas, though the formal zone selection mechanism is introduced in detail in a later chapter.

Overall, the operational setting illustrates that the problem is a daily service planning task under operational constraints rather than a long-term static assignment problem. The interaction between limited team capacity, geographically dispersed stations, fixed working hours, and dynamically emerging disturbances creates a planning environment in which route design and real-time adaptation are equally important.

2.2 Environment and Disturbances

A defining feature of the planning problem is the presence of uncertainty. The maintenance teams operate in an environment in which service requirements are not fully known at the start of the day. New faults may occur while teams are already executing their planned routes, and the operational system therefore evolves over time. For this reason, a purely static routing approach is insufficient. Instead, the planning logic must account for the fact that the set of required service actions can change during execution.

In this setting, disturbances are treated as exogenous operational events. They occur stochastically and may emerge at different charging stations at different times of the day. Once a disturbance occurs, it immediately creates an additional service requirement that competes for limited maintenance capacity. This means that service planning must continuously balance already scheduled preventive work with newly arriving corrective tasks. The problem is therefore not only about finding efficient routes, but also about deciding how to reallocate attention when the system state changes unexpectedly.

The route-based model distinguishes between two types of disturbances. Type-1 faults are simple failures that can be resolved through standard repair actions with a fixed service time. Type-2 faults are more complex and require disassembly, a depot roundtrip, and reassembly, resulting in a significantly longer and station-dependent service time. This distinction is operationally important because the insertion of a short Type-1 job and the insertion of a long Type-2 repair do not have the same impact on the remainder of the day.

Uncertainty is not limited to the occurrence of disturbances. Travel conditions also vary over time and are subject to uncertainty. Travel times depend on the time of day due to varying traffic conditions, and their actual realization is stochastic. This means that geographically similar routes may still differ substantially in execution quality depending on temporal traffic patterns and local congestion effects.

Another relevant feature of the environment is that unresolved disturbance jobs do not disappear at the end of the day. If a newly occurring fault cannot be inserted feasibly into the remaining route of either team, it is transferred as a carryover job to the following day. These carryover disturbances are then prioritized in the next daily planning cycle. This mechanism introduces an intertemporal dependency into the operational system, since today's capacity shortage affects tomorrow's workload. In practical terms, the planning problem therefore includes not only immediate routing consequences but also backlog management over multiple days.

2.3 Assumptions and Cost Model

To make the planning problem analytically tractable, the operational system is represented through a number of simplifying assumptions. First, not every real-world process is modeled in full physical detail. Some side processes associated with complex repairs are incorporated through aggregate service-time components instead of explicit intermediate route legs. In the current formulation, a Type-2 fault includes additional time components for disassembly, depot-related handling, and reassembly, but the team remains conceptually assigned to the affected station rather than physically traveling to and from the depot as a separate route segment. This simplification reduces model complexity while preserving the main operational time impact of such events.

Second, service requirements are represented as operational jobs rather than as fully persistent technical condition states. Open regular maintenance tasks, new faults, and carryover disturbances are treated as actionable job sets. In addition, each charging station

is associated with a maintenance-related status through the indicator of whether periodic maintenance has been completed and through the variable days since maintenance, which later drives failure risk. However, the binary distinction between “healthy” and “faulty” infrastructure is not carried as a separate persistent state variable in the operational problem description. This abstraction supports a cleaner transition into the route-based MDP in Chapter 4.

The economic evaluation of the planning problem is based on three cost components. The first component consists of labor costs for the maintenance teams, set to 35 € per hour. This figure is derived from an average gross wage of approximately 21 € per hour for maintenance technicians in Bavaria (jobvector GmbH, 2024), scaled up by employer-side social security contributions of roughly 20–25 % and an overhead allowance for provisions and downtime, resulting in a total employer cost of approximately 35 € per hour. The second component consists of travel costs, set to 0.30 € per kilometer, corresponding to the standard statutory mileage reimbursement rate for business travel in Germany (Bundesministerium der Finanzen, 2024). The third component consists of downtime-related costs, set to 0.50 € per kilowatt-hour, approximating the market price for public AC charging in Germany as reported by the BDEW (Bundesverband der Energie- und Wasserwirtschaft, 2024). This rate reflects the opportunity cost of lost charging capacity per kilowatt-hour that could not be delivered due to an unresolved fault. Together, these three components reflect that good operational performance depends not only on efficient routing, but also on timely fault resolution.

The cost structure creates a clear trade-off. A route with low travel distance is not necessarily desirable if it leaves important faults unserved for too long. Conversely, a strongly service-oriented plan that aggressively inserts every disturbance may lead to high labor and travel expenditure or may jeopardize the timely return to the depot. The relevant operational objective is therefore not the minimization of a single physical quantity such as distance, but the minimization of total expected cost across labor use, mobility effort, and service delay consequences.

2.4 Motivating Example

To illustrate why both the initial plan and the replanning decision require more than routing efficiency, consider the simplified scenarios shown in Figures 2 and 3. The scenario involves four stations, A, B, C, and D, and a central depot. Station B has not been serviced in 365 days and operates at 80 kW, placing it at maximum failure risk. Station A was serviced only 30 days ago and carries only 11 kW, making it still non-critical. Stations C and D have intermediate maintenance states with 100 and 200 days since last service respectively.

Figure 2 contrasts two approaches to constructing the initial route. The left panel shows a routing-efficiency approach, in which stations are ordered primarily by geographic convenience, resulting in Depot $\rightarrow$ B $\rightarrow$ A $\rightarrow$ C $\rightarrow$ D. While B is correctly visited first, A is scheduled second despite its low urgency, consuming an early slot that could have served a more critical station. The right panel shows a maintenance-urgency approach, in which

2 Problem Definition




Figure 2: Two initial plan strategies for the same set of stations. Left: routing-efficiency approach visits A second despite low urgency. Right: maintenance-urgency approach defers A to last, ensuring critical stations are covered first.

stations are ordered by their criticality, deferring A to the last position. The route becomes Depot → B → D → C → A. The key advantage becomes apparent when a disruption arrives later in the day and forces a drop, since in the right panel, the critical stations B, D, and C have already been serviced, so A is the natural candidate for removal with negligible consequences. In the left panel, a disruption arriving after A might force the drop of C or D, stations with considerably higher maintenance urgency.

This shows that even the initial plan is a decision with future consequences, since the order in which stations are visited determines which ones remain at risk if capacity is unexpectedly consumed by a disruption later in the day.

Figure 3 illustrates the replanning situation. Starting from the planned route, disturbance E occurs during execution and is inserted as the first stop after the depot. One station must be dropped to keep the route feasible. The left panel shows a purely distance-driven choice, in which B is dropped because it causes the largest routing detour, resulting in E → A → C → D. B remains unserved, its days since maintenance has already reached 365 days and its failure risk remains at its maximum, a cost that will likely materialize the following day. The right panel shows a state-aware decision in which A is dropped instead, as its low days since maintenance and low rated power make a one-day delay negligible. The route becomes E → B → D → C, B is visited, and its maintenance counter resets to zero, averting the failure risk entirely.

Both examples point to the same insight, namely that good maintenance routing requires knowledge of each station's state, not just geographic distance. The illustrative scenario uses days since maintenance and rated power as the two most prominent state variables, but in practice additional station-specific characteristics such as station age or geographic

2 Problem Definition




Figure 3: Two replanning decisions under the same incoming disruption E. Left: dropping B (largest detour) leaves the most critical station unserved. Right: dropping A (recently serviced, low power) allows B to be serviced and averts failure risk.

context may also influence the urgency of a service visit. A purely distance-driven approach ignores this information and can systematically produce worse outcomes. The central idea of the solution approaches developed in this thesis is therefore to assign an explicit value to each station based on its current maintenance state, and to use this value as a scoring criterion in both the initial plan construction and the replanning decision. This shifts the problem from a purely geometric routing task to a state-dependent optimization problem, in which the expected future cost of leaving a station unserved is weighed against the immediate routing detour required to serve it.

Chapter 3 reviews the relevant literature on maintenance optimization, stochastic vehicle routing, and approximate dynamic programming, providing the theoretical foundation for the solution approach. Building on both, Chapter 4 translates the operational setting into a route-based Markov decision process and connects that formulation to an approximate dynamic programming solution approach.

3 Literature Review

The problem outlined in Chapter 2 sits at the intersection of three research streams, namely maintenance optimization, stochastic vehicle routing, and ADP. This chapter reviews each of these streams in turn and shows how they collectively provide the theoretical basis for the solution approach developed in this thesis.

3.1 Maintenance Optimization in Infrastructure Systems

Maintenance optimization in infrastructure systems addresses the challenge of planning, scheduling, and executing maintenance activities to ensure system availability, reliability, and cost-efficiency over the lifecycle of technical assets. Unlike manufacturing or logistics contexts, infrastructure maintenance often involves geographically distributed assets, limited maintenance resources, and significant consequences of system failures (De Jonge and Scarf, 2020).

The following subsections trace this progression, from the core objectives and trade-offs of maintenance optimization, through the dominant maintenance strategies and their limits, to the specific complications that arise when technicians must travel between geographically distributed sites, the operational structure directly relevant to this thesis.

3.1.1 Concept and Objectives of Maintenance Optimization

Maintenance optimization refers to the systematic determination of when and what to maintain in order to achieve specific objectives such as minimizing costs, maximizing availability, or balancing both (De Jonge and Scarf, 2020). The fundamental trade-off lies between the cost of maintenance activities and the cost of system failures or downtime (Dekker, 1996).

Early work by Dekker (1996) identified six barriers preventing the application of maintenance optimization models in practice, namely that models are difficult to understand, many papers are written only for mathematical purposes, maintenance activities consist of many different aspects, optimization is not always necessary, models often focus on the wrong type of maintenance, and there is a gap between theory and operational reality. Despite these challenges, maintenance optimization remains strategically important because continuous developments of technical systems and increasing reliance on equipment have resulted in growing importance of effective maintenance activities (De Jonge and Scarf, 2020).

In the context of charging infrastructure, maintenance optimization must address not only individual station failures but also fleet-level dispatching, resource allocation, and timing decisions under uncertainty (De Jonge and Scarf, 2020). The objective is typically to minimize total expected costs (maintenance plus failure costs) while maintaining a required level of system availability (De Jonge and Scarf, 2020).

3.1.2 Maintenance Strategies: Preventive, Reactive, Condition-Based

Maintenance strategies can be broadly classified into three categories, namely reactive (corrective) maintenance, preventive (planned) maintenance, and condition-based (predictive) maintenance (Mołęda et al., 2023; De Jonge and Scarf, 2020).

Reactive maintenance (also called corrective or run-to-failure maintenance) is performed only after a failure occurs (Mołęda et al., 2023). This strategy is suitable for non-critical assets where downtime has minimal impact on operations. However, for critical infrastructure such as charging stations, reactive maintenance can lead to significant customer dissatisfaction, safety risks, and higher repair costs due to unanticipated failures (De Jonge and Scarf, 2020).

Preventive maintenance is performed according to a fixed schedule or usage-based interval, regardless of the actual condition of the asset (Mołęda et al., 2023; Alaswad and Xiang, 2017). Preventive maintenance aims to prevent failures before they occur by replacing or servicing components at predetermined times. This approach is well-established and widely applied, for example in vehicle maintenance or periodic inspections of technical systems (Dekker, 1996). However, preventive maintenance can lead to over-maintenance (replacing components that are still functional) or under-maintenance (failures occurring before the scheduled interval) (De Jonge and Scarf, 2020).

Condition-based maintenance (CBM) performs maintenance based on the monitored condition of the asset, for example sensor data or error rates, and is often extended to predictive maintenance using machine learning or digital twins (Alaswad and Xiang, 2017). While CBM is well-studied in the literature, it requires continuous condition monitoring infrastructure that is not assumed in this thesis.

In practice, many organizations use a hybrid strategy, combining preventive maintenance for critical components with reactive maintenance for non-critical ones and condition-based maintenance for assets where monitoring is available (De Jonge and Scarf, 2020). The charging infrastructure context falls naturally into this hybrid territory, since routine inspections address predictable deterioration, while spontaneous hardware failures demand immediate reactive response, a combination that motivates the dynamic replanning logic developed in Chapter 5.

3.1.3 Uncertainty, Availability, and Cost as Central Decision Variables

Maintenance decisions are inherently made under uncertainty. Key uncertain parameters include failure times, service times, travel times, availability of spare parts, and arrival of new maintenance requests (De Jonge and Scarf, 2020). These three dimensions, availability, cost, and uncertainty, jointly define the scope of maintenance optimization:

- **Availability** measures the proportion of time that a system is operational and able to fulfill its function (Alaswad and Xiang, 2017). For charging infrastructure, failed stations directly reduce customer access and revenue, making availability a primary

performance indicator.

- **Cost** encompasses maintenance expenditures (labor, materials, travel), failure costs (downtime, customer dissatisfaction, penalties), and inventory costs (spare parts, tools) (De Jonge and Scarf, 2020; Dekker, 1996). The optimization objective is typically to minimize total expected cost over a planning horizon, subject to availability constraints (Alaswad and Xiang, 2017).
- **Uncertainty** affects both the timing and the magnitude of maintenance consequences. A failure may occur earlier or later than expected, service may take longer than planned, or travel to a site may be delayed by traffic. Decision-making therefore follows a sequential process in which uncertainty gradually resolves and decisions are revised in response (De Jonge and Scarf, 2020).

In maintenance optimization, uncertainty is typically modeled through stochastic deterioration processes, probability distributions for failure times, or scenario-based approaches, with the choice depending on data availability, system complexity, and computational requirements (Alaswad and Xiang, 2017).

3.1.4 Transition to Mobile Maintenance and Service Interventions

The maintenance problems discussed in Sections 3.1.1–3.1.3 typically focus on when and what to maintain at individual assets. However, in many practical settings, maintenance is performed by mobile technicians who must travel between geographically distributed sites, which introduces routing and scheduling decisions (De Jonge and Scarf, 2020).

For charging infrastructure maintenance, this means that the optimization problem is not only about determining maintenance intervals or failure thresholds but also about dispatching technicians, planning routes, and allocating limited resources (vehicles, spare parts, time) across multiple stations (De Jonge and Scarf, 2020).

This transition from single-asset maintenance optimization to mobile maintenance service planning introduces additional complexity:

- **Spatial dependency:** Maintenance decisions for one station affect the availability of resources for other stations (De Jonge and Scarf, 2020).
- **Temporal dependency:** The order in which stations are visited affects total travel time and service availability (Alaswad and Xiang, 2017).
- **Uncertainty coupling:** Failures, travel times, and service times are often stochastic and interdependent (De Jonge and Scarf, 2020).

The literature on maintenance optimization has traditionally focused on single-asset or multi-asset models without explicit routing (De Jonge and Scarf, 2020). More recent work begins to integrate maintenance planning with vehicle routing, especially in the context of stochastic and dynamic vehicle routing problems (SDVRP) (De Jonge and Scarf, 2020). This integration is precisely the focus of the following Sections 3.2 and 3.3, which cover stochastic routing and approximate dynamic programming for dynamic service planning.

3.1.5 EV Charging Infrastructure as a Maintenance Context

The maintenance of electric vehicle charging infrastructure constitutes an emerging and practically relevant application within the broader field of infrastructure maintenance optimization. As EV adoption accelerates globally, the scale and geographic dispersion of charging networks have grown substantially, making systematic maintenance planning a critical operational challenge for charging spot operators (al., 2024).

EV charging stations exhibit two distinct operational requirements that together define the maintenance workload. First, charging piles require frequent inspections to identify potential safety hazards such as overheating, communication failures, or display errors. While individual inspections are relatively simple, the sheer number of stations makes inspection routing a non-trivial planning problem. Second, a substantial share of stations are subject to outright failures, as empirical data from Chinese EV charging operators indicate that more than ten percent of deployed charging spots contain faulty charging piles at any given time, necessitating corrective maintenance by qualified engineers (al., 2024). Taken together, operations and maintenance costs account for approximately 30–50 % of the total cost of running a charging spot network, making efficient maintenance routing a significant economic lever (al., 2024).

Despite this practical relevance, the academic literature on EV charging station maintenance routing remains thin. The closest related work is al. (2024), who propose a bi-objective routing optimization model for the large-scale inspection-maintenance of EV charging spots in Xi'an, China, covering 1,281 stations. Their multi-stage approach applies K-Means clustering to reduce problem scale, estimates cluster-level costs via a greedy TSP heuristic, and solves the reduced model with an adaptive NSGA-II to obtain a Pareto frontier balancing total cost and composite customer satisfaction loss. This is, to the best of the author's knowledge, one of the first papers to explicitly address routing optimization for EV charging station maintenance.

However, the model of al. (2024) is static and deterministic, since all maintenance and inspection tasks are planned once at the beginning of the horizon, with no mechanism for intraday disruptions, stochastic fault arrivals, or replanning. In practice, charging stations fail unpredictably during the working day, and maintenance teams must revise their routes in response. This dynamic dimension, recurring daily over a multi-month operational horizon, with state-dependent task priorities that change as stations age between visits, is precisely what the present thesis addresses. The contribution of this work therefore lies not in replacing static routing models, but in extending the planning paradigm toward a dynamic, state-aware, and intertemporally consistent maintenance policy for EV charging infrastructure.

3.2 Stochastic Vehicle Routing Problems and Dynamic Service Planning

Efficient service planning in real-world environments is rarely a one-time optimization problem. Instead, it involves a sequence of decisions under uncertainty, where new in-

formation becomes available over time and must be incorporated into the planning process. This is particularly true for maintenance operations on charging infrastructure, where breakdowns, service requests, and priority changes can emerge unpredictably during the day.

The literature on Vehicle Routing Problems (VRP) has evolved from static, deterministic formulations toward dynamic and stochastic models that better reflect such operational realities (Ritzinger, Puchinger, and Hartl, 2016). This section provides the theoretical foundation for understanding how stochastic and dynamic routing problems are modeled and solved.

3.2.1 Basic Concepts of VRP and DVRP

The classical Vehicle Routing Problem (VRP) involves planning optimal routes for a fleet of vehicles to serve a set of customers under various operational constraints (Pillac et al., 2013). Common constraints include vehicle capacity and service time windows (Hildebrandt, Thomas, and Ulmer, 2023). In the static VRP, all relevant information is known at the beginning and does not change during planning (Ritzinger, Puchinger, and Hartl, 2016). The dynamic Vehicle Routing Problem (DVRP) extends this basic model by incorporating information that emerges over time, for example new requests, changing travel times, or additional service needs (Abbatecola, Fanti, and Ukovich, 2016). The stochastic Vehicle Routing Problem (SVRP) further augments this dynamism by explicitly modeling uncertainty, i.e., random or only partially known events (Akhmetbek and Mustakhov, 2023). For the literature, it is important that “dynamic” and “stochastic” often occur together but are not identical, since a problem can be dynamic without being strongly stochastic, or stochastic without being re-planned in real time (Ritzinger, Puchinger, and Hartl, 2016). The survey literature shows that research has increasingly moved from the static VRP toward dynamic and stochastic variants because real logistics and service processes are rarely fully planable (Mardešić et al., 2023). For this thesis, this transition is central because maintenance tasks for charging infrastructure cannot be planned completely deterministically either (Ulmer, Goodson, Mattfeld, and Hennig, 2019).

3.2.2 Uncertainties in Routing: Requests, Service Times, Travel Times, Capacities

A typical uncertainty aspect in the Stochastic Dynamic Vehicle Routing Problem (SDVRP) is the arrival of new customer requests or service requests that occur only during the tour (Zou and Dessouky, 2018). While Ulmer, Goodson, Mattfeld, and Hennig (2019) focus on general service requests like roadside assistance, this thesis applies this logic to maintenance orders and spontaneous failures at charging infrastructure. Another uncertainty factor is service times, which are hardly predictable exactly in practice because repairs may take different amounts of time or technical problems may occur on site (Akhmetbek and Mustakhov, 2023). Travel times can also be stochastic, for example due to traffic, weather, roadworks, or time-dependent speeds (Mardešić et al., 2023). Additionally, capacities can be uncertain, for example when vehicles carry only limited tools, spare parts, charging energy, or

personnel resources (Basso et al., 2022). The literature usually distinguishes several types of uncertainty, including stochastic demand, stochastic travel times, and stochastic arrival times (Ulmer, Mattfeld, and Köster, 2018). For modeling, it is important which uncertainties are known *ex ante* and which are observed only online during the tour (Ulmer, Goodson, Mattfeld, and Hennig, 2019). Particularly relevant for this work is the fact that uncertainty affects not only the optimal sequence of stops but also whether an intervention is worthwhile at all or whether one should wait for further information (Ulmer, Goodson, Mattfeld, and Hennig, 2019). It is precisely at this point that the transition from classical routing to dynamic service planning becomes visible (Ulmer, Goodson, Mattfeld, and Hennig, 2019).

3.2.3 Modeling Approaches: MDP, Multi-Stage Stochastic Programming, Look-ahead

Many works in the reviewed literature model the SDVRP as an MDP because decisions must be made sequentially under uncertainty (Ulmer, Goodson, Mattfeld, and Hennig, 2019). In an MDP, the problem is described by states, actions, transition probabilities, and an objective function (Ulmer, Goodson, Mattfeld, and Hennig, 2019). The state can contain vehicle positions, remaining time, open requests, priorities, or current loads (Ulmer, Goodson, Mattfeld, and Hennig, 2019). An action is then, for example, the choice of the next intervention site, waiting for new information, or the decision whether to accept a job directly (Ulmer, Goodson, Mattfeld, and Hennig, 2019). The major advantage of this modeling is that it explicitly represents the dynamics of the problem and incorporates future uncertainty into the decision (Ulmer, Goodson, Mattfeld, and Hennig, 2019). However, the literature also emphasizes that an exact MDP quickly becomes very large because the number of possible states and actions explodes (Ulmer, Goodson, Mattfeld, and Hennig, 2019). Therefore, approximate approaches such as ADP or look-ahead-based procedures are frequently used in practice (Ulmer, Goodson, Mattfeld, and Hennig, 2019). Alternatively, some works employ multi-stage stochastic programming models, where multiple decision stages with uncertainty realizations are optimized (Ritzinger, Puchinger, and Hartl, 2016). These models are conceptually elegant but often difficult to solve for large, realistic instances (Ritzinger, Puchinger, and Hartl, 2016). Look-ahead procedures, by contrast, consider only a limited planning horizon and attempt to make the short-term best decision while taking some future scenarios into account (Zou and Dessouky, 2018). For this thesis, this modeling family is particularly relevant because it lies closer to real maintenance and dispatching processes than purely static tour planning (Ulmer, Goodson, Mattfeld, and Hennig, 2019). The central insight from the literature is that the more uncertainty and real-time dynamics are embedded in the problem, the more one needs MDP- or look-ahead-based formulations instead of classical routing models (Ulmer, Goodson, Mattfeld, and Hennig, 2019).

3.2.4 Solution Families: Exact Methods, Heuristics, Metaheuristics

The literature on SDVRP shows several solution families that differ primarily in accuracy, scalability, and online capability (Ritzinger, Puchinger, and Hartl, 2016). Exact methods attempt to solve the problem mathematically optimally, for example via Mixed-Integer Programming (MIP), dynamic programming, or exact search procedures (Pillac et al., 2013).

These approaches are useful for small instances but quickly reach limits with dynamic, stochastic, and large-scale problems (Ritzinger, Puchinger, and Hartl, 2016). Therefore, heuristic procedures dominate the SDVRP literature, i.e., rules or search procedures that deliver good but not necessarily optimal solutions (Mardešić et al., 2023). Typical heuristics include myopic rules, greedy dispatching, and insertion heuristics, which are particularly attractive in dynamic settings because they allow fast online decisions (Pillac et al., 2013). Metaheuristics such as Variable Neighborhood Search or Tabu Search can improve solution quality but tend to handle future uncertainty less explicitly than ADP-based approaches, making them less suitable as the primary method in anticipative maintenance planning (Ulmer, Goodson, Mattfeld, and Hennig, 2019). For this thesis, greedy insertion serves as the operational routing primitive within all developed policies. Its adoption is motivated by three properties that are particularly valuable in a dynamic maintenance context. First, it runs in polynomial time, enabling real-time route revisions whenever disruptions arrive during the day. Second, it integrates naturally with priority-based station scoring, allowing maintenance urgency to influence the insertion order directly. Third, it provides a transparent and interpretable baseline on which the ADP layer can build without obscuring the contribution of the anticipative component (Pillac et al., 2013). The anticipative element is then introduced through the ADP framework described in Section 3.3.

3.2.5 Route-based MDPs and Their Relevance for This Thesis

Route-based MDPs are particularly important because they model the routing problem at the level of the route rather than only at the level of individual states or customer requests. Ulmer, Goodson, Mattfeld, and Thomas (2020) formally introduce this framework, defining route-based MDPs as a generalization of conventional MDPs in which the action space, state space, and reward structure are augmented to operate on routing plans rather than individual node selections. The advantage lies in the fact that the route itself is understood as a decision unit, and not only individual stops are considered in isolation. This allows anticipative decisions, i.e., decisions that do not only react to the current state but consider expected future events (Ulmer, 2020). Particularly relevant are the works on Approximate Dynamic Programming, offline-online ADP, and horizontal combinations of on-line and offline methods (Ulmer, 2020). These works show that route-based MDPs can be well combined with ADP strategies because the value function or rollout logic can approximate future developments (Ulmer, 2020). For this thesis, this is a strong argument because maintenance interventions also do not consist only of individual decisions but of a sequence of interdependent route and intervention decisions (Ulmer, Goodson, Mattfeld, and Hennig, 2019). Particularly fitting is the idea that failures or service requirements occur during the day and therefore are not fully known in advance (Ulmer, Goodson, Mattfeld, and Hennig, 2019). The route-based perspective allows exactly this kind of dynamic, sequential, and resource-aware planning (Ulmer, 2020). Route-based MDPs therefore serve as the methodological hinge between the SDVRP literature reviewed in this section and the ADP framework introduced in Section 3.3. They provide the formal structure within which anticipative decision-making for charging station maintenance can be both modeled and

approximated (Ulmer, Goodson, Mattfeld, and Hennig, 2019).

3.2.6 Limits of Classical SDVRP Models for Maintenance Tasks

Classical SDVRP models often originate from urban logistics, last-mile delivery, or courier-service contexts (Mardešić et al., 2023). These applications are structurally similar but not fully transferable to maintenance logistics (Abbatecola, Fanti, and Ukovich, 2016). In maintenance scenarios, different decision priorities arise, such as technical criticality, failure consequences, recommissioning, and personnel qualification (Mardešić et al., 2023). Furthermore, maintenance interventions are often less standardized than deliveries because diagnosis, repair, and feedback can vary greatly (Abbatecola, Fanti, and Ukovich, 2016). Many SDVRP models treat requests as customer orders, whereas for this thesis, failure or maintenance orders are more relevant (Ulmer, Goodson, Mattfeld, and Hennig, 2019). This means that not only the spatial sequence but also technical prioritization and temporal urgency play a role (Ulmer, Goodson, Mattfeld, and Hennig, 2019). Another difference is that for charging infrastructure maintenance, return trips, spare parts, technical rework, and multi-visit scenarios can occur frequently (Abbatecola, Fanti, and Ukovich, 2016). Many literature models abstract these aspects strongly and are therefore only conditionally realistic (Ritzinger, Puchinger, and Hartl, 2016). The strength of the SDVRP literature thus lies less in a direct template for this problem but rather in transferable decision principles (Mardešić et al., 2023). There are strong methods for dynamic routing problems, but the specific maintenance context of charging infrastructure remains underrepresented. This gap motivates the approach taken in this thesis, adopting the structural logic of route-based SDVRPs while adapting the priority rules, cost model, and replanning logic to the requirements of technical maintenance operations (Ulmer, Goodson, Mattfeld, and Hennig, 2019).

3.2.7 Spatial Clustering as a Pre-Structuring Mechanism

A recurring approach in EV charging maintenance routing is the use of spatial clustering to decompose a large station network into manageable service zones before applying route optimization. al. (2024) apply K-Means clustering to a network of 1,281 EV charging stations in Xi'an, China, where clustering is used to reduce the combinatorial scale of the problem, since maintenance stations are retained individually while inspection spots are grouped into clusters, and the resulting reduced problem is then solved with an adaptive NSGA-II metaheuristic to minimize total cost and customer satisfaction loss simultaneously. Their multi-stage approach demonstrates that geographic pre-structuring is a practically necessary step when the network scale would otherwise make direct optimization intractable. Yenigun, Baydogmus, and Yildiz (2025) adopt a similar spatial decomposition for 100 stations in Istanbul, grouping stations via K-Means and subsequently optimizing inter-cluster routing with a genetic algorithm, achieving a 35 % reduction in total travel distance compared to a sequential baseline. Both works share the core idea with this thesis that K-Means clustering over station coordinates provides a natural and computationally efficient way to organize service zones. However, a key difference is that both papers treat clustering as

a one-time, static pre-processing step applied to the full station network before routing begins. In this thesis, by contrast, zone membership is fixed but zone *selection* is a daily decision, made at the start of each working day, when the system evaluates all open zones and assigns teams to those with the highest current priority, taking into account backlog, station criticality, and team positions. This repeated, state-dependent zone selection introduces a dynamic layer that static clustering approaches do not address and constitutes one of the structural contributions of the thesis.

3.3 Approximate Dynamic Programming for Dynamic Routing Problems

The dynamic and stochastic nature of real-world service planning requires decision frameworks that can handle sequential choices under uncertainty without prohibitive computational costs. Exact Dynamic Programming (DP) provides a theoretically optimal foundation but quickly becomes intractable for realistic routing problems. ADP addresses this limitation by approximating future costs or values rather than computing them exactly, thereby enabling scalable, anticipative decision-making in operational settings. This section introduces the core concepts, methods, and relevance of ADP for dynamic routing.

3.3.1 Motivation: Why Exact Dynamic Programming Does Not Scale

DP is theoretically elegant for routing problems with uncertainty but often practically intractable because the state space explodes rapidly. Even with moderate fleet sizes, multiple open requests, and dynamic events, an enormous combinatorial structure arises (Ulmer, 2017). The literature consistently shows that exact DP serves as a reference but is too computationally intensive for real-world SDVRP instances (Ulmer, Goodson, Mattfeld, and Hennig, 2019). This is precisely where ADP begins. Instead of computing the exact optimal solution, ADP approximates future costs or future utility well enough to make good decisions. While Ulmer, Mattfeld, and Köster (2018) focus on the time-sensitive nature of parcel pickup services, their logic regarding the need for fast and robust decisions is directly applicable to the maintenance context. In both domains, the dispatcher cannot wait for mathematical perfection but must react immediately to stochastic events, such as spontaneous breakdowns, to ensure operational capability. The ADP literature thereby justifies the transition from theoretical optimality to operational decision-making power (Mardešić et al., 2023). Particularly the Ulmer works show that routing problems with stochastic requests can realistically only be solved online with approximate procedures (Ulmer, Goodson, Mattfeld, and Hennig, 2019). Exact DP remains valuable as a conceptual benchmark, but ADP provides the practical framework for real-time service planning under uncertainty.

3.3.2 Basic Idea of ADP: States, Actions, Value Functions

ADP treats the problem as a sequential decision under uncertainty. A state describes the current planning status, for example vehicle positions, remaining time, open orders, the route traveled so far, or available resources (Ulmer, Goodson, Mattfeld, and Hennig, 2019). An action is a concrete decision in the current state, for example which order to serve next

or whether to wait for new information (Ulmer, 2017). The value function describes how good a state or a state-action combination is in the long run. The central point is that future costs or returns are not computed exactly but approximated (Ulmer, 2017). This approximation can happen via linear functions, features, simulations, or data-driven approximations (Ulmer, 2020). That is precisely why ADP is more flexible than exact Dynamic Programming approaches. Ulmer (2017) and Ulmer, Goodson, Mattfeld, and Hennig (2019) apply this logic to evaluate routing decisions with expected future quality. For this work, this way of thinking is useful because maintenance interventions also start from a current state, and the quality of a decision only shows itself over the future. In mathematical terms, ADP approximates the Bellman equation:

$$V_t(s_t) \approx \mathbb{E} \left[\sum_{k=t}^T \gamma^{k-t} C_k(s_k, a_k) \mid s_t \right]$$

where $V_t(s_t)$ is the approximated value of state s_t at time t , s_k and a_k are the state and action at stage k , C_k is the resulting cost, and γ is a discount factor (Ulmer, 2017). Instead of computing this expectation exactly, ADP uses approximation functions learned offline or updated online.

3.3.3 Important ADP Approaches: Rollout, VFA, and Offline-Online ADP

Rollout is one of the most practically important ADP heuristics. A base policy is used as a reference, and for possible actions an anticipative future simulation is performed (Ulmer, Goodson, Mattfeld, and Hennig, 2019). Rollout is particularly strong when one must decide quickly online but still wants to account for the future at least approximately (Ulmer, Goodson, Mattfeld, and Hennig, 2019). VFA replaces the exact value function with an approximated function. The approximation can be based on features, partitionings, or historical data (Ulmer, 2020). Across the reviewed literature, VFA is repeatedly identified as a dominant ADP component, especially in the survey papers and in the offline-online approaches (Mardešić et al., 2023). Offline-Online ADP combines an offline learned or pre-computed value function with an online executed decision rule (Ulmer, Goodson, Mattfeld, and Hennig, 2019). This is particularly attractive because computational effort is shifted into the preparation phase while remaining responsive online (Ulmer, Goodson, Mattfeld, and Hennig, 2019). As Ulmer, Goodson, Mattfeld, and Hennig (2019, p. 185) note:

The offline-online policy enhances the anticipation of the VFA policy, yielding spatial and temporal anticipation of requests and routing developments. The combination of VFA with rollout algorithms demonstrates the potential benefit of using offline and online methods in tandem as a hybrid ADP procedure, making possible higher-quality policies with reduced computational requirements for real-time decision making.

A related concept worth noting is the post-decision state, the system state immediately after an action is taken but before random information is revealed (Ulmer, 2017). This separation of deterministic action effects from subsequent stochasticity is conceptually useful,

and the rollout-based policies developed in Chapter 5 implicitly rely on it, since each Monte Carlo scenario begins from the post-decision state after a drop or insertion choice, before new faults are drawn. For this thesis, Rollout, VFA, and Offline-Online ADP are the three central paradigms. Their combination of learned value weights applied offline and rollout evaluation applied online is precisely the structure realized in the solution policies of Chapter 5 (Ulmer, Goodson, Mattfeld, and Hennig, 2019). In this thesis, the rollout layer serves directly as the value approximation mechanism and is therefore referred to as the VFA component, consistent with the ADP literature's view that rollout itself constitutes a form of value function approximation (Ulmer, 2017).

3.3.4 Distinction from Reinforcement Learning

Reinforcement Learning (RL) and ADP share the sequential decision framework but differ in origin, since RL emphasizes learning from interaction data, while ADP operates from the Operations Research tradition using explicit model-based structure (Mardešić et al., 2023). For complex route-based problems, ADP tends to be more robust because it separates search and evaluation explicitly via value functions or rollout logic, whereas many RL approaches scale poorly or fail to capture routing combinatorics (Hildebrandt, Thomas, and Ulmer, 2023). This thesis follows the ADP paradigm, using simulation-based rollout and learned value weights rather than end-to-end RL policy training.

3.3.5 Relevance of ADP for Dynamic Maintenance Decisions

ADP is particularly suitable for maintenance decisions because maintenance is typically not decided once but sequentially and under uncertainty (Ulmer, Goodson, Mattfeld, and Hennig, 2019). A defective charger, an unplanned failure, or a changing priority are exactly the type of events that ADP can address (Ulmer, Mattfeld, and Köster, 2018). The future does not need to be known exactly, since it is sufficient to approximate it plausibly and align the decision with it (Ulmer, Goodson, Mattfeld, and Hennig, 2019). Offline-Online ADP is particularly attractive when maintenance rules can be trained or precomputed in advance and then applied quickly in operation (Ulmer, Goodson, Mattfeld, and Hennig, 2019). Rollout is sensible when on-site decisions must be made between several interventions or locations and the consequences of options should be approximately simulated (Zou and Dessouky, 2018). VFA is useful when one wants to build a robust evaluation of states from experience data or simulated scenarios (Ulmer, 2020). Post-decision States are useful because they cleanly separate the effect of an action from the subsequent entry of uncertainty (Anuar et al., 2022). Overall, the literature shows that ADP is not only very well transferable to classical delivery logistics but also to dynamic technical service processes (Mardešić et al., 2023). For this thesis, the core message is therefore that ADP provides a scalable, anticipative, and practice-oriented decision framework for maintenance planning under uncertainty (Ulmer, Goodson, Mattfeld, and Hennig, 2019). In the context of charging infrastructure maintenance, this means that ADP can handle breakdowns that emerge during the day, prioritize interventions based on technical criticality, and allocate limited resources such as spare parts or technician time efficiently (Ulmer, Mattfeld, and Köster, 2018). The offline-online

combination allows maintenance policies to be learned from historical failure data or simulations and then deployed in real-time dispatching systems (Ulmer, Goodson, Mattfeld, and Hennig, 2019).

4 Theoretical Framework

This chapter develops the formal framework used to represent the maintenance routing problem as a route-based MDP (Ulmer, Goodson, Mattfeld, and Thomas, 2020), in which actions correspond to full route plans rather than individual node selections. Building on the operational setting introduced in Chapter 2, the objective is now to translate the planning problem into a mathematical structure that captures sequential decision-making under uncertainty. Particular emphasis is placed on the interaction between route construction, stochastic fault arrivals, time-dependent travel conditions, and limited daily service capacity. The chapter further explains why a state-action-value perspective is required and why ADP constitutes an appropriate solution methodology for the resulting problem. Table 6 in the Appendix summarizes the notation used throughout this chapter.

4.1 Markov Decision Process

The maintenance of electric vehicle charging infrastructure is inherently dynamic. At the start of the working day, not all information relevant for operational planning is yet known, because additional disturbances may occur during execution and traffic conditions may affect route feasibility. Consequently, the planning problem cannot be represented adequately by a purely static routing model. Instead, it is formulated as an MDP, in which decisions are made sequentially and the system evolves stochastically over time.

A route-based MDP formulation is especially suitable in this context because route decisions have consequences beyond their immediate effect. Assigning a team to one area of the network influences the remaining time budget, the accessibility of future service jobs, and the ability to react to newly emerging disturbances. The purpose of the formulation is therefore not merely to choose the next stop, but to evaluate route decisions in light of their effect on the future development of the system.

4.1.1 Sets, indices, and temporal structure

The maintenance network consists of a set of charging stations $I = \{1, \dots, 397\}$ and a set of service teams $K = \{1, 2\}$. The depot is indexed as node 0, and the overall node set is given by $V = \{0\} \cup I$. Travel times between locations are denoted by d_{ij} for all $i, j \in V$ and may vary by time of day. In addition, the charging stations are partitioned into zones $Z = \{1, \dots, N_Z\}$, with $N_Z = 178$, which serve as an operational pre-structuring mechanism for route generation.

The decision process is embedded in a daily planning horizon. The working day begins at 8:00 and planning decisions are evaluated on an hourly grid $H = \{8, 9, \dots, 16\}$. Within this

structure, two types of decision points are distinguished. First, an initial planning decision is made at the beginning of the day, where a full route plan is constructed for both teams. Second, replanning decisions may occur during the day whenever new disturbances arise. This distinction between an initial plan and event-triggered route revisions reflects the operational logic of a maintenance system in which execution and planning must remain closely coupled.

Although each day constitutes a self-contained planning instance within this finite horizon, successive days are not independent. The days-since-maintenance variable dsm_d^i and the carryover tasks from unresolved faults create an explicit inter-day coupling. Servicing a station on day d lowers its fault probability on day $d + 1$, while leaving stations unattended raises the carryover backlog and increases expected future fault rates. The total planning objective therefore accumulates daily costs over a full operational year of D days. Within a single day, the Q^* -function captures the downstream consequences of a decision for the remainder of that day, while across days, these consequences propagate through the DSM update and the carryover mechanism.

The lunch break between 12:00 and 13:00 constitutes a deterministic component of the transition function rather than an explicit decision variable. It is inserted after the first stop whose scheduled departure falls at or after noon, shifting all subsequent arrival and departure times by one hour. Since the break applies uniformly and unconditionally, it requires no dedicated state variable or action representation, and is instead absorbed directly into the transition operator $P(s_{h+1} \mid s_h, a_h)$ as a fixed timing rule.

4.1.2 State representation

Let the system state at decision epoch $h \in H$ be denoted by

$$s_h = (\tau_h, X_h, Y_h, J_h, \Pi_h).$$

This representation is designed so that the state contains all information required to make the next decision. In that sense, it satisfies the Markov property, since the probability distribution of the next state depends on the current state and the chosen action, but not on earlier system history beyond what is already encoded in s_h .

The first component, τ_h , captures the current time within the workday and is measured in minutes since 8:00. The second component, X_h , describes the state of the service teams. For each team $k \in K$, this includes its current location v_h^k , its activity status σ_h^k , and the remaining time budget b_h^k available before the route must return to the depot, defined as

$$b_h^k = T_{\max} - \tau_h - d_{v_h^k, 0},$$

where $T_{\max}$ denotes the total length of the working day in minutes and $d_{v_h^k, 0}$ is the travel time from the team's current location back to the depot. The activity status distinguishes whether a team is currently servicing a station, travelling, returning to the depot, or inactive. This part of the state is essential because route feasibility depends critically on the remaining usable time and on the team's current physical position.

The third component, Y_h , describes the station-related maintenance status. For each station $i \in I$, the state contains a maintenance indicator m_h^i , which records whether the required periodic maintenance has already been completed, and a variable dsm_h^i , representing the number of days since the last maintenance intervention. The latter variable plays a central role because it influences the probability of future disturbances. Importantly, the technical condition of a charging station is not represented here as a separate persistent binary or multi-level health state. Instead, disturbances are modeled as stochastic events that are immediately transformed into service requests.

The fourth component, J_h , captures the set of open maintenance jobs. It consists of three subsets, namely regular maintenance jobs W_h^{reg} , newly arrived fault jobs W_h^{fault} , and carry-over fault jobs W_h^{carry} transferred from the previous day. This distinction is operationally relevant because the planning problem must account simultaneously for preventive tasks, newly occurring corrective tasks, and accumulated backlog. The set of open regular maintenance jobs is directly determined by the maintenance indicators: $W_h^{\text{reg}} = \{i \in I : m_h^i = 0\}$.

Finally, Π_h stores the planned routes of both service teams. For each team, the route plan contains a node sequence and the associated schedule of arrival and departure times. In a route-based formulation, including the route plan as part of the state is particularly useful because actions do not merely select an immediate successor node. Rather, they modify or reconstruct larger route segments and thereby influence the remaining structure of the entire working day.

4.1.3 Action space

The action space $A(s_h)$ depends on the type of decision point. At the start of the day, the decision consists of constructing an initial route plan for both teams. This process is divided into two stages. In the first stage, zones containing open stations are evaluated and filtered. The scoring can be based on classical geographic criteria such as average distance to the depot and zone area, or on a value-based priority measure reflecting maintenance urgency and station-specific relevance. Each team is then assigned a start zone, and stations are added to the candidate set by loading further zones in descending score order until a predefined maximum number of stations is reached.

In the second stage, the actual route plan is constructed. Based on the candidate sets and the carryover jobs inherited from the previous day, a feasible route for each team is generated by a greedy cheapest-insertion heuristic. Stations that cannot be accommodated within the remaining working time are dropped from the current day and remain open. This route construction procedure yields the full-day action taken at the initial decision point.

At replanning points during the day, the action space has a different interpretation. Newly arriving disturbance jobs are considered sequentially and inserted into the existing route of the team for which the additional insertion cost is smallest. The insertion cost of placing fault j at position pos in the route is

$$c_{\text{ins}}(j, \text{pos}) = \frac{\Delta t_{\text{extra}}}{60} \cdot c_L + \Delta \text{km} \cdot c_F + \Delta \tau_j^{\text{wait}} \cdot p_j \cdot c_{\text{dt}},$$

where Δt_{extra} is the additional travel time caused by the detour, Δkm the corresponding extra distance, and $\Delta \tau_j^{\text{wait}}$ the waiting time from fault report until arrival. Since this downtime term grows with the waiting time until arrival, earlier insertion positions are almost always cheaper, meaning disruptions are in practice served as the next stop immediately after they are reported. If no feasible insertion exists because the remaining route time is insufficient, the disturbance is not ignored but transferred as a carryover job to the following day.

All actions must satisfy three feasibility conditions. First, each route must be temporally feasible, meaning that the planned return to the depot must occur within the allowable time horizon. Second, each team is required to end its route at the depot. Third, each station may be served by at most one team at any given time. These constraints ensure that every action corresponds to an operationally executable route plan. The chosen action at decision epoch h is denoted $a_h \in A(s_h)$.

4.1.4 Transition dynamics

The transition function $P(s_{h+1} \mid s_h, a_h)$ describes how the system evolves between two consecutive decision points. During that interval, teams execute the planned route, travel to stations, perform maintenance or repair tasks, and update the completion status of serviced stations. Simultaneously, new disturbances may occur and travel times may deviate from their nominal values. As a result, the transition dynamics are stochastic even if the chosen action is fixed.

The route progression is determined by the planned arrival and departure schedule, but actual execution is affected by uncertain travel conditions. The formulation therefore allows for stochastic travel times, for example by using time-dependent lognormal travel-time realizations. This captures the operational reality that the duration of a trip depends not only on spatial distance but also on the time of day and traffic conditions.

Service completion updates the maintenance-related station state. A completed routine maintenance task sets the maintenance indicator to one and resets the days-since-maintenance variable to zero. A type-1 fault requires a standard repair with fixed duration $s_{\text{typ1}} = 60$ min and likewise resets the maintenance age. A type-2 fault entails a more extensive service duration that incorporates additional handling and depot-related effort through an aggregated time term. In the current route-based approximation, this depot-related part is not modeled as a separate physical route leg but is embedded directly in the service time. The aggregated service time for a type-2 fault at station i is therefore given by

$$s_{\text{typ2}}(i) = s_{\text{dismount}} + d_{i,0} + s_{\text{handling}} + d_{0,i} + s_{\text{remount}},$$

with $s_{\text{dismount}} = 30$ min, $s_{\text{handling}} = 5$ min, and $s_{\text{remount}} = 30$ min. This modeling choice preserves the operational impact on route feasibility without introducing additional state and routing complexity.

At the end of each day d , where d indexes calendar days and $h \in H$ indexes decision epochs

within a day, the days-since-maintenance variable is updated according to

$$dsm_{d+1}^i = \begin{cases} 0 & \text{if station } i \text{ received routine maintenance or fault repair on day } d, \\ dsm_d^i + 1 & \text{otherwise.} \end{cases}$$

Disturbance generation is modeled stochastically and depends on station-specific baseline probabilities, maintenance age, and station-level scaling factors. For each station $i \in I$ and fault type $\ell \in \{1, 2\}$, the per-hour fault probability at decision epoch h is

$$p_\ell^i(h) = p_{\ell, \text{base}} \cdot f(dsm_h^i) \cdot \eta^i,$$

where $p_{\ell, \text{base}}$ is the baseline fault rate, η^i is a station-level scaling factor, and $f(\cdot)$ is a recovery curve that captures how the fault probability increases with maintenance age:

$$f(t) = \lambda + (1 - \lambda) \cdot \frac{\min(t, \tau_{\text{rec}})}{\tau_{\text{rec}}}.$$

Here, $\lambda \in (0, 1)$ is the residual fault factor immediately after a maintenance visit and τ_{rec} is the number of days until full fault probability is restored. Faults that cannot be scheduled feasibly in the current day are carried over and become part of the next day's initial workload. Through this mechanism, unresolved work propagates over time and creates a dynamic coupling between successive planning periods.

4.1.5 Reward function

The optimization objective is formulated as a cost minimization problem. In MDP notation, the reward is therefore defined as the negative of the total cost incurred by an action and its resulting system evolution:

$$R = -(C_{\text{op}} + C_{\text{dt}}).$$

The total cost consists of an operational component and a downtime component:

$$C_{\text{op}} = n_{\text{active}} \cdot T_{\text{shift}} \cdot c_L + \sum_{(i,j) \in \text{trips}} \text{km}_{ij} \cdot c_F,$$

$$C_{\text{dt}} = \sum_{j \in \text{faults}} \Delta\tau_j^{\text{wait}} \cdot p_j \cdot c_{\text{dt}}.$$

The operational cost C_{op} includes labor and travel expenditures. Labor costs depend on the number of active teams n_{active} and the effective shift duration T_{shift} , while travel costs depend on the kilometers km_{ij} traveled between successive route segments. In the present formulation, the labor cost parameter is $c_L = 35$ € per hour and the travel cost parameter is $c_F = 0.30$ € per kilometer.

The downtime-related cost C_{dt} is accounted for on a per-day basis, summing over all unresolved fault jobs. Here p_j denotes the nominal charging power of station j in kilowatts and $c_{\text{dt}} = 0.50$ € per kilowatt-hour the downtime unit rate. The effective waiting time $\Delta\tau_j^{\text{wait}}$

depends on when the fault is resolved. For a fault that cannot be inserted into any team's route on the day it arrives, $\Delta\tau_j^{\text{wait}}$ equals the remaining workday duration, while for a carryover fault resolved on the following day, it equals the elapsed workday time from the start of operations until service completion. A fault that remains unresolved over multiple days accumulates daily downtime costs on each affected day independently. This structure ensures that the objective does not reward short routes per se, but rather seeks an economically balanced compromise between efficient deployment and timely restoration of charging availability.

Overall, the reward function reflects the central trade-off of the maintenance problem. A route that minimizes travel distance may still perform poorly if it postpones high-impact repairs, while an aggressively reactive policy may generate excessive operating cost. The chosen reward structure therefore evaluates actions in terms of their total contribution to operational efficiency and service quality.

Together, the state space S , action space $A(s_h)$, transition function P , and reward R constitute the route-based MDP at the operational, hourly level. For the value analysis in the following section, this fine-grained process is aggregated to the day level: the decision attributed to day d is the initial route plan, while the replanning actions taken at the within-day decision epochs follow the fixed reactive insertion and drop rules of the respective policy and are therefore absorbed, together with stochastic fault arrivals and travel-time realizations, into the day-level transition kernel. The planning objective is to find an optimal policy $\pi^* : S \rightarrow A$ that minimizes the expected total cost over the full planning year. As formally defined in Section 4.2, this is expressed through the optimal state-action-value function Q^* :

$$\pi^*(s_d) = \arg \min_{a_d \in A(s_d)} Q^*(s_d, a_d).$$

4.2 State-Action-Value Structure

While the state representation describes the current condition of the system, it is not sufficient by itself to evaluate decision quality. In a dynamic maintenance routing problem, the attractiveness of a decision depends not only on its immediate cost, but also on how it shapes the future development of the system. This is the reason for adopting a state-action-value perspective.

The optimal state-action-value function $Q^*(s, a)$ assigns to each feasible action a in state s the minimum expected total downstream cost obtained when that action is chosen and the process then continues optimally. Since the total objective accumulates daily costs over the planning year, the Bellman equation is formulated at the day level, where d indexes calendar days:

$$Q^*(s_d, a_d) = \mathbb{E} \left[C(s_d, a_d) + \min_{a_{d+1} \in A(s_{d+1})} Q^*(s_{d+1}, a_{d+1}) \mid s_d, a_d \right],$$

where $C(s_d, a_d) = C_{\text{op}} + C_{\text{dt}}$ is the total cost incurred on day d and the expectation is taken

over the stochastic transition $s_{d+1} \sim P(\cdot \mid s_d, a_d)$, which embeds all within-day disruption handling and stochastic travel-time realizations. The formulation is undiscounted ($\gamma = 1$), meaning that future costs are weighted equally to immediate costs, a deliberate choice reflecting that postponing a repair is equally costly regardless of when it occurs within the planning horizon. Conceptually, Q^* combines two elements, namely the immediate daily cost of the action and the minimum expected future cost achievable from the resulting state. The latter component is crucial because route decisions alter the set of feasible responses to later disturbances.

This distinction is especially important in a route-based MDP. In a node-based formulation, an action might simply correspond to selecting the next customer or the next destination. In the present case, the action a_d corresponds to the full initial route plan constructed at the start of day d . Within-day disruption handling, the sequence of insertion and drop decisions triggered by arriving faults, is not part of a_d but is embedded in the stochastic transition $P(s_{d+1} \mid s_d, a_d)$, which captures all within-day dynamics including replanning, stochastic travel times, and fault arrivals. This day-level view is consistent with the hourly formulation of Section 4.1: the replanning actions defined there on the hourly grid are not free decision variables at the day level, but are executed by the fixed reactive decision rules of the policy under evaluation, so conditioning on these rules renders the day-level transition kernel well-defined. Consequently, the quality of an initial plan cannot be assessed only by looking at its planned route cost. It must instead be judged in terms of how well it positions the teams to handle the disruptions that will unfold during execution.

The practical meaning of this is straightforward. Inserting a station into the route consumes time, changes the expected depot return time, affects accessibility of later jobs, and may either preserve or reduce the ability to react to future disturbances. A locally attractive decision may therefore be globally poor if it fills the route with low-priority stations that must later be dropped to accommodate high-priority faults. Conversely, a locally more expensive insertion can be preferable when it maintains route adaptability and reduces expected future service delay.

The Q -structure provides the conceptual foundation for evaluating such trade-offs. It translates the routing problem from a purely myopic assignment logic into a sequential decision problem in which future consequences are explicitly recognized. In this way, the state-action-value perspective forms the bridge between the route-based MDP formulation and the approximation strategy introduced in the next section.

4.3 ADP Foundations and the Offline-Online Approach

Although the route-based MDP provides an appropriate conceptual representation of the maintenance problem, solving it exactly is computationally infeasible. The dimensionality of the state space is very large because it combines temporal information, two vehicle states, station-specific maintenance information, multiple categories of open jobs, and the current route plan. The action space is equally large, since route construction and route modification decisions can take many feasible forms. In addition, the transition process is

stochastic due to fault arrivals and uncertain travel times.

For these reasons, the problem is addressed through approximate dynamic programming (ADP). The central idea of ADP is not to compute the exact optimal value function for every possible state, which would be impractical, but to approximate the value of future decisions in a computationally tractable way. This approximation can be based on simulation, parameterized value estimates, structured heuristics, or combinations thereof. The purpose is to obtain a policy that is sufficiently accurate to support high-quality operational decisions while remaining fast enough for day-to-day deployment.

In the present context, ADP is particularly attractive because the decision problem unfolds in two distinct modes, namely planning at the beginning of the day and replanning during execution. A suitable framework must therefore support both anticipatory route construction and rapid reaction to new disturbances. This naturally leads to an offline-online structure.

In the offline phase, the decision logic is prepared before actual daily operation. This can include calibration of scoring parameters, design of value surrogates, generation of representative simulation scenarios, and estimation of how route characteristics translate into future cost consequences. The offline phase is therefore the stage at which structural knowledge about the maintenance system is embedded into the approximation framework.

In the online phase, the system observes the current state and must select a feasible action under tight time constraints. The precomputed or learned information from the offline stage is then used to guide route construction or route revision. The degree of online computation varies across policies. Simpler approaches rely entirely on precomputed parameters for fast online decisions, while more advanced approaches additionally perform targeted online simulation, using the precomputed base policy, to refine those decisions. In both cases, the offline preparation is what makes the approach practical for operational use.

The route-based logic of the present model fits this ADP structure particularly well. The zonal pre-selection, candidate generation, and greedy insertion rules should not be interpreted merely as ad hoc heuristics. Within the ADP framework, they can be viewed as operational policy components that exploit structural information about the value of different routing choices. Their role is to restrict the search space intelligently and to approximate the effect of decisions on future system performance.

Taken together, ADP provides the methodological layer that makes the route-based MDP operationally usable. The route-based MDP defines the decision problem rigorously, the state-action-value structure clarifies what has to be approximated, and the offline-online ADP scheme supplies a practical mechanism for producing fast and robust route decisions under uncertainty. Chapter 5 translates this framework into concrete policies, ranging from a myopic baseline that ignores future costs entirely, through cost function approximations that embed learned station values into the routing score, to a rollout-based lookahead that approximates the value function via short-horizon Monte Carlo simulation.

5 Solution Methodology

Building on the MDP framework and ADP foundations developed in Chapter 4, this chapter presents the solution methodology developed for the maintenance planning problem of electric charging stations. The methodological framework is designed to handle a dynamic and stochastic environment in which new failures may occur during the day and existing routes must be adjusted accordingly. Since the problem is too complex to solve exactly at the operational level, the thesis relies on a hierarchy of heuristic policies and ADP-based approaches.

The methodology starts with a common spatial decision layer that determines which zones of the network are considered first. On top of this shared structure, different decision policies are applied. The simplest policy is the Myopic baseline, which scores stations purely by travel time without any economic weighting or anticipation of future costs. Myopic+ extends this baseline by incorporating a stronger economic prioritization. CFA introduces a learned station value that captures the future cost of postponing service. DB extends this cost-based structure by introducing a state-dependent routing parameter. VFA extends this hierarchy by adding an online rollout layer that replaces the greedy drop decision with a simulation-based lookahead.

The purpose of this chapter is to describe these policies in a consistent way and to make clear how they differ in terms of decision logic, prioritization, and replanning behavior. This structure also allows for a fair comparison of the methods in the computational experiments.

Figure 4 summarizes this hierarchy. Each policy to the right extends its predecessor by adding one further layer of economic reasoning, from a purely distance-driven baseline (Myopic) to a state-dependent balance between value and routing efficiency (DB), while the VFA rollout layer sits above all four and can wrap any of them without altering their internal logic.

5.1 Shared Zone Selection Logic

The zone selection mechanism forms a preliminary decision layer that acts before the actual routing and disturbance handling logic. Its purpose is to determine which parts of the service area are addressed first on a given day. All policies share the same underlying zone structure, based on a K-Means decomposition of the station network, but differ in how zones are scored. Zone selection is evaluated as an independent experimental dimension, so each policy is tested under both a centrality-based and a value-based scoring mode, which allows the effect of zone selection to be isolated from the routing and replanning logic.

In the centrality-based mode, zones are scored by their geographic centrality within the




Figure 4: Hierarchy of decision policies developed in this thesis. All four base policies share the same zone-selection layer (Section 5.1) and differ only in their routing score and drop score, shown in each box. The VFA layer (Section 5.6) is not a fifth base policy but an optional rollout wrapper that can be placed on top of any of the four.

network. The score is derived from each zone's mean distance to all other zone centroids:

$$\text{score}(z) = 1 - \frac{\bar{d}_z - \min_z \bar{d}_z}{\max_z \bar{d}_z - \min_z \bar{d}_z},$$

where $\bar{d}_z$ denotes the mean distance from zone z 's centroid to the centroids of all other zones. Zones that are close to many other zones receive a high score and are processed first. This criterion is policy-agnostic and serves as the baseline for zone prioritization.

In the value-based mode, the zone score is instead determined by the expected maintenance relevance of the stations it contains. Specifically, the score is the sum of a policy-specific station value over all open stations in the zone. The station value function differs by policy. Myopic and Myopic+ use a power-weighted urgency measure $v(i) = p_i \cdot \rho(\text{dsm}_i)$, where p_i is the rated power and $\rho(\text{dsm}_i)$ is the recovery curve evaluated at the days since last maintenance, while CFA and DB use the learned cost approximation $\tilde{C}$. This allows zone selection to reflect not only geographic structure but also the anticipated cost consequences of postponing service.

At the start of each day, all open zones are scored and ranked according to the active mode. From the top 20 zones by score, each team is then assigned the zone that lies geographically closest to its current position, so the highest-ranked zone is not necessarily selected if it is far from both teams. From this start zone, each team's candidate set is filled with additional stations by loading further zones in descending score order until a maximum of 20 stations per team is reached. Team 1 fills its candidate set first, and stations it has claimed are no longer available to Team 2. Both teams then build their routes independently from their respective candidate sets using the policy-specific routing logic described in the following sections.

When a disruption arrives during the day, it is evaluated for insertion into both teams' current routes simultaneously, and the team and position with the lowest insertion cost are selected. Drop decisions are made independently per team. Stations dropped from one team's route are never transferred to the other but return to the open maintenance pool for future days.

5.2 Myopic

The Myopic policy serves as the baseline heuristic for all subsequent methods. It scores stations purely by travel time without any economic weighting or lookahead, and does not incorporate any form of learned cost approximation. This makes it a fast, transparent, and easily interpretable reference point for evaluating the added value of the more advanced policies.

Initial Plan

At the beginning of the day, after zone selection and candidate set generation, the initial route is constructed by a greedy routing procedure. Starting from the depot, the next station is always chosen as the one with the highest score among all remaining candidates, where the score is defined as the inverse of the travel time from the current position:

$$\text{score}(i, \text{cur}) = \frac{1}{\max(1, d_{\text{cur},i})}.$$

This rule prioritizes nearby stations and produces a compact, travel-time-efficient route without any economic weighting. Carryover tasks from previous days are always inserted first, using a nearest-neighbor rule, before any routine tasks are scheduled. This ensures that previously unresolved faults are addressed as early as possible in the new day. Routine stations that cannot be accommodated within the remaining workday are dropped and remain open for the following day.

Disruption Handling

When new disruptions arrive during the day, each disruption is evaluated for insertion into the current route using a greedy cheapest-insertion procedure. For every combination of

team and insertion position, the additional cost is computed as:

$$c_{\text{ins}} = \frac{\Delta t_i + s_i}{60} \cdot c_L + \max(0, \Delta \ell_i) \cdot c_F + \frac{\max(0, t_i^{\text{arr}} + s_i - r_i)}{60} \cdot p_i \cdot c_{\text{dt}},$$

where the detour time and detour distance are defined as

$$\Delta t_i = d_{\text{prev},i} + d_{i,\text{next}} - d_{\text{prev},\text{next}}, \quad \Delta \ell_i = \ell_{\text{prev},i} + \ell_{i,\text{next}} - \ell_{\text{prev},\text{next}},$$

s_i is the service time of the disruption in minutes, $c_L = 35$ €/h the labor rate, $c_F = 0.30$ €/km the fuel rate, t_i^{arr} the arrival time at the disrupted station in minutes since 8:00, r_i the fault report time in minutes since 8:00, p_i the rated power in kW, and $c_{\text{dt}} = 0.50$ €/kWh the downtime cost rate. Here d_{ij} denotes travel time in minutes between nodes i and j as defined in Chapter 4, and ℓ_{ij} the corresponding travel distance in kilometers. Disruptions are processed in ascending order of their cheapest insertion cost, so the least expensive disruption is inserted first.

If no feasible direct insertion exists, a drop-and-insert fallback is applied. All remaining routine stops are ranked by their detour time Δt_i , defined as the travel time saved by removing stop i from the current route. Stops are dropped in ascending order of the score $1/\max(1, \Delta t_i)$, meaning the stop that saves the largest detour is removed first. Drops are applied iteratively until the disruption can be feasibly inserted. Dropped routine stops return to the open maintenance pool and are reconsidered in future days through normal zone selection. If no combination of drops creates sufficient capacity, the disruption itself becomes a carryover task for the following day.

5.3 Myopic+

Myopic+ shares the same greedy architecture as Myopic but replaces the purely travel-time-based scoring with a power-weighted score. The central idea is that stations with higher rated power should be visited earlier in the day, since a failure at such a station incurs higher downtime costs. This economic weighting is applied consistently in both the initial plan and the drop decision during disruption handling.

Initial Plan

Carryover tasks are inserted first using nearest-neighbor, identical to Myopic. For routine tasks, the next station is selected by the highest score among remaining candidates, where the score now weights travel time by rated power:

$$\text{score}(i, \text{cur}) = \frac{p_i}{\max(1, d_{\text{cur},i})}.$$

A station with high rated power p_i and short travel time $d_{\text{cur},i}$ receives the highest priority. Compared to Myopic, a more distant high-power station can therefore be preferred over a nearby low-power station if the power advantage outweighs the travel penalty.

Disruption Handling

The insertion cost formula for evaluating disruption candidates is identical to Myopic, since the same full economic cost over detour time, distance, and downtime is computed for every (team, position) pair, and disruptions are processed in ascending order of insertion cost.

The key difference lies in the drop score. When no direct insertion is feasible, remaining routine stops are ranked by:

$$\text{drop_score}(i) = \frac{p_i}{\max(1, \Delta t_i)},$$

where Δt_i is the detour time saved by removing stop i . Stops are dropped in ascending order of this score, so low-power stations with large detours are removed first, while high-power stations or those with small detours are protected. This ensures that economically important stations remain in the route whenever capacity is tight. Dropped stops are carried over to the following day, and if no combination of drops resolves the capacity shortage, the disruption itself becomes a carryover.

5.4 Cost Function Approximation

CFA extends the greedy decision logic by introducing a learned approximation of the future cost of postponing a station. Instead of scoring stations purely by proximity or rated power, CFA estimates how costly it would be to leave a station unserved and uses this estimate as a prioritization signal throughout the planning day. The approximation is linear in a small set of station-specific features and trained offline, while all operational decisions remain fully greedy.

Cost Approximation and Training

CFA represents the cost of leaving station i unserved through a linear value function over four features:

$$\phi(i) = [p_i, a_i, \rho(\text{dsm}_i), \bar{d}_i]^\top,$$

where p_i is the rated power, a_i the station age in years, $\rho(\text{dsm}_i)$ the recovery curve evaluated at the days since last maintenance, and $\bar{d}_i$ the mean distance from station i to all other stations, precomputed statically at initialization. The recovery curve captures the temporal failure risk, since it is lowest immediately after maintenance and rises back toward its maximum as time passes. Before entering the model, all features are standardized with training means and standard deviations, and the value function is:

$$\tilde{C}(i, \text{dsm}_i) = \theta^\top \phi_{\text{scaled}}(i, \text{dsm}_i).$$

The weight vector $\theta \in \mathbb{R}^4$ is estimated offline via Ridge regression on contrastive training labels. Each label is the discounted cost difference between two simulated suffix trajectories for station i , one in which the station is served and one in which it is dropped. This makes $\tilde{C}$ a direct approximation of the marginal value of servicing a station rather than a

fit to aggregate rollout costs. The standardization parameters μ and σ are stored together with θ and loaded at inference time.

Initial Plan

Carryover tasks are inserted first using nearest-neighbor, as in Myopic and Myopic+. For routine tasks, the station selection score combines the learned value with travel time:

$$\text{score}(i, \text{dsm}_i, \text{cur}) = \frac{\tilde{C}(i, \text{dsm}_i) + s}{\max(1, d_{\text{cur},i})},$$

where $s = \max(0, -\min_i \tilde{C}) + 1$ is a shift that ensures all scores are strictly positive, preventing negative $\tilde{C}$ values from reversing the travel-time penalty. When value-based zone selection is active, CFA uses $\tilde{C}$ as the station value function passed to the zone selector, as described in Section 5.1.

Disruption Handling

The insertion cost formula is identical to Myopic, since full economic costs over detour time, distance, and downtime are evaluated for each (team, position) pair, and disruptions are processed in ascending order of insertion cost. When no direct insertion is feasible, remaining routine stops are ranked by a drop score that trades off future value against the detour savings from removal:

$$\text{drop_score}(i, \text{dsm}_i) = \tilde{C}(i, \text{dsm}_i) - \frac{c_L}{60} \cdot \Delta t_i,$$

where $c_L/60 = 35/60$ €/min is the per-minute labor rate and Δt_i is the detour time saved by removing stop i . Stops are dropped in ascending order, so stations with low future value and large detour savings are removed first. Dropped stops become carryovers, and if no combination of drops creates sufficient capacity, the disruption itself is carried over to the following day.

5.5 Dynamic Balance

DB builds directly on CFA and extends it by introducing a state-dependent balancing parameter δ . The purpose of this policy is to dynamically trade off the economic importance of a station, captured by the approximated value $\tilde{C}(i)$, against routing efficiency, captured by the distance term in the greedy selection rules. In contrast to CFA, where this trade-off is fixed, DB allows the emphasis to shift with the state of the system.

Balance Concept and Parameter

Following the Dynamic Balance concept of Stein et al. (2025), routing quality should not be determined by a single static priority rule. Instead, the decision logic should balance two competing objectives, namely the servicing of economically important stations and

the maintenance of spatial efficiency in the route. This trade-off is particularly relevant in a dynamic maintenance setting, where the best choice early in the planning horizon may not be the best choice as the horizon closes.

The balance parameter $\delta \in \{0.5, 0.9\}$ controls the relative strength of the distance term in the greedy score. The special case $\delta = 0.5$ yields a routing score with a linear distance penalty, which is identical to CFA. Smaller values of δ reduce the importance of distance and shift the focus more strongly toward the approximated value $\tilde{C}(i)$. Larger values increase the distance penalty and thereby place greater emphasis on routing compactness. The parameter is determined once before the start of each day and remains constant throughout the day. In the experimental setup of this thesis, δ follows a simple binary threshold, remaining at $\delta = 0.5$, identical to CFA, as long as at least 130 stations remain unserved, and switching to $\delta = 0.9$ once fewer than 130 stations remain. In practice, this means the policy behaves like CFA throughout most of the planning horizon and shifts toward more compact routing only in the final stage.

Initial Plan

DB inherits all core components from CFA without modification: the feature vector $\phi(i)$, the learned value function $\tilde{C}(i, \text{dsm}_i)$, the shift s , the carryover nearest-neighbor insertion, and the value-based zone selection. The only change is in the routing score, where the distance exponent is now controlled by δ :

$$\text{score}(i, \text{dsm}_i, \text{cur}) = \frac{\tilde{C}(i, \text{dsm}_i) + s}{\max(1, d_{\text{cur},i})^{2\delta}}.$$

When $\delta = 0.5$, the exponent equals one and the formula reduces exactly to the CFA score. For $\delta < 0.5$, distance is penalized less than linearly, so high-value stations can be reached at a greater travel cost. For $\delta > 0.5$, the distance penalty grows super-linearly, favoring compact routing over value-based prioritization.

Disruption Handling

The insertion cost formula is identical to CFA and Myopic. When no direct insertion is feasible, routine stops are removed according to a drop score that mirrors the same δ -controlled balance:

$$\text{drop_score}(i, \text{dsm}_i) = \tilde{C}(i, \text{dsm}_i) - 2\delta \cdot \frac{c_L}{60} \cdot \Delta t_i,$$

where Δt_i is the detour time saved by removing stop i and $c_L/60 = 35/60$ €/min. The factor 2δ scales the weight of the detour savings relative to the future value. At $\delta = 0.5$ this reduces to the CFA drop score. For larger δ , stops with high detour savings are dropped more aggressively, improving feasibility at the cost of dropping potentially valuable stations.

5.6 Value Function Approximation

VFA is not an independent planning algorithm in this thesis, but an online rollout layer that is placed on top of an existing base heuristic. The base policy remains fully intact and continues to generate the operational plan. VFA only intervenes at selected decision points, where a rollout simulation indicates that an alternative action yields a better expected future outcome. In this sense, VFA is an offline-online hybrid strategy in which the heuristic structure is defined in advance and the improvement is obtained online through simulation-based evaluation.

Offline-Online Architecture

The VFA framework is based on a strict separation between a static offline component and a dynamic online decision component. In the offline part, the chosen base policy is fixed and prepared for use in the rollout layer. In the online part, the current decision state is evaluated by simulating alternative actions and comparing their expected future consequences over a finite horizon. The central idea is not to replace the base policy, but to complement it with a future-oriented evaluation mechanism. This separation allows the method to combine computational tractability with a stronger anticipation of future events.

The offline-online structure is important because it keeps the operational logic simple while still enabling more informed decisions at critical points. The offline policy provides the reference behavior, whereas the online rollout approximates the value of different actions by simulating their downstream effects. This is consistent with the general ADP principle that a good heuristic base policy can be improved by a local lookahead layer. In the context of this thesis, this architecture makes it possible to retain the strengths of the previously defined policies while adding a more anticipatory decision mechanism.

Offline Phase: Base Policy

In the offline phase, a base policy is selected as the foundation for the rollout layer. Depending on the experimental setup, this base policy may be Myopic, Myopic+, CFA, or DB. The chosen policy remains unchanged during the rollout process and continues to handle the standard operational planning. The role of the offline phase is therefore to provide a stable and fully specified heuristic structure on which the rollout can build.

The base policy is accessed through a unified interface, which allows the VFA layer to remain independent of the concrete heuristic used underneath. This is important because it makes the rollout mechanism modular and transferable across different policy classes. For the initial plan, the base policy remains fully responsible in the standard configuration. In the standard setting, the offline phase does not perform any additional optimization itself, but rather defines the behavior that will later be evaluated and, if necessary, overridden by the online layer.

This means that VFA does not construct routes from scratch. Instead, it uses the base policy as the default operational rule and improves it only when the rollout suggests a better local

alternative. The offline phase is therefore best understood as the structural backbone of the method. It guarantees that the system always has a feasible and interpretable default behavior, even if the rollout layer is not activated or cannot find a superior candidate.

Online Phase: Rollout

The actual value approximation takes place in the online phase through rollout simulations. VFA is activated when a newly reported disturbance cannot be inserted directly and a drop decision must be made. The base policy ranks all remaining routine stops by its drop score and identifies the top- k candidates ($k = 3$ in all experiments). For each candidate, the system constructs a post-decision state, the system state after that drop has been applied but before future uncertainty materializes, and simulates the remaining planning horizon using the base policy over a horizon of 40 days and 25 stochastic scenarios per candidate.

To ensure a fair comparison across candidates, all candidates are evaluated under the same set of stochastic scenarios, generated using common random numbers. This reduces sampling noise and makes the cost estimates statistically comparable. A candidate overrides the base-policy choice only if it yields a lower expected cost in at least 70 % of the simulated scenarios and produces a positive average gain. The 70 % threshold was chosen deliberately, since lower thresholds such as 60 % admitted too many marginal overrides driven by sampling noise rather than genuine cost differences, leading to higher variance in outcomes. The 70 % filter ensures that only decisions with clear and consistent evidence of improvement are adopted. If no candidate clears this threshold, or if the rollout computation exceeds a predefined time budget, the system falls back to the greedy base-policy decision.

Rollout is applied exclusively at drop decision points and not during initial plan construction. Applying rollout to the initial plan was explored in preliminary experiments but did not yield consistent improvements, likely because the combinatorial structure of full-day route construction makes isolated candidate comparisons less informative than at the level of individual drop decisions.

Overall, VFA represents the highest level in the methodological hierarchy of this thesis. It combines the robustness of a heuristic base policy with the anticipatory power of simulation-based lookahead. Compared with CFA and Dynamic Balance, it is more strongly future-oriented because it does not merely approximate station relevance, but directly evaluates the expected downstream cost of alternative drop decisions. At the same time, it remains operationally feasible because rollout is applied only at drop decision points rather than throughout the entire planning process.

6 Experimental Setup

This chapter specifies the simulation environment and data basis used to compare the maintenance policies. The setup is designed to capture both the operational structure of the

charging-station maintenance problem and the stochastic effects that influence routing and downtime.

6.1 Simulation Design

The policy evaluation is based on a discrete-time simulation of the daily maintenance operations described in Chapter 2 and formally specified in Chapter 4, run over a horizon of up to 365 days. Because both failure generation and travel times are stochastic, a single run does not produce reliable performance estimates, so the simulation is repeated in Monte Carlo experiments. All stations are initialized at maximum days-since-maintenance, representing a fully deteriorated starting state. This is a conservative choice that avoids any warm-up bias in favor of recently serviced stations and ensures that failure risk is present from the first day of simulation.

Each policy is evaluated over 500 independent runs using seeds 1–500. Two random streams are maintained separately throughout each run, one governing stochastic failure generation and the other governing travel-time perturbations. Keeping these streams independent avoids unintended correlation between the two uncertainty sources. Crucially, the same seed produces the same disruption sequence regardless of which policy is being evaluated. Pairwise comparisons between policies are therefore matched by seed, which eliminates run-to-run noise as a confounding factor and substantially increases statistical power.




Figure 5: Execution timeline of a single simulated day for both teams (CFA, value-based zone selection, seed 1, day 1), reconstructed from the actual simulation log. Both teams start their routine route at 8:00; disruptions (Type 1 and Type 2) are inserted mid-route via cheapest insertion, and the shared 60-minute lunch break is fixed at 12:00–13:00 and interrupts the active service stop if one is ongoing. Of the 20 routine tasks planned for the day, 10 could not be completed and were carried over to day 2.

Figure 5 illustrates this mechanism concretely for a single simulated day. Both teams begin at the depot at 8:00 and initially follow a route of routine stops, and as disruptions are reported throughout the day, they are inserted directly into the ongoing route, and low-

priority routine stops that no longer fit are dropped and carried over to the following day, as described in Section 5.2.

6.2 Routing and Travel Times

Travel times are constructed from two separate sources. Base travel times, free of any traffic influence, are obtained from the Open Source Routing Machine (OSRM) running on a local OpenStreetMap extract for Bavaria. A complete 398×398 matrix covering the depot and all 397 stations is precomputed once and reused throughout the simulation.

Time-dependent congestion is added on top of these base times using the Google Distance Matrix API. A set of relevant traffic nodes in the Würzburg road network is queried hourly from 8:00 to 17:00 for a representative weekday. For each node pair and each hour, the delay is computed as the difference between the queried travel time and a 3:00 a.m. baseline, which is assumed to be free of congestion. These hourly delay values are then matched to simulation routes via their geometric footprint, so any route whose path passes through a queried node segment receives the corresponding delay addition. Routes that do not intersect a queried segment receive the mean delay observed across matched routes for that hour. The result is a separate traffic matrix for each of the nine working hours, capturing the morning and afternoon rush-hour peaks as well as the lower congestion during midday.

To represent short-term travel-time variability beyond the deterministic hourly averages, each simulated leg is additionally perturbed by a log-normal factor. The coefficient of variation is set to 0.20 during the morning peak (8:00), decreases to 0.10 during midday, and rises again to 0.15–0.20 in the afternoon. These values are plausible assumptions rather than estimates from empirical traffic measurements.

6.3 Data Basis

The simulation is based on a real data set of publicly available charging infrastructure in the city of Würzburg, obtained from the municipal open data portal (Stadt Würzburg, 2024). The portal provides station-level records for the entire public charging network in Würzburg, including geographic coordinates, charging power, number of charging points, commissioning date, and the commercial operator. The data set contains 397 operationally active charging stations, consisting of both normal charging stations and fast-charging stations. These stations are operated by 83 distinct commercial providers. The municipal utility Stadtwerke Würzburg AG, a WVV subsidiary, is the single largest operator with 145 stations (36.5 %), while the remaining 252 stations (63.5 %) are distributed across 82 other providers. Station age is inferred from the commissioning date relative to the start of the simulation period.

The depot is placed at a WVV facility in Sanderau, chosen as a geographically central starting point for both maintenance teams. This choice of depot location does not imply that WVV operates the entire station network modeled here, but rather that the two maintenance teams represent a hypothetical, city-wide coordinating entity responsible for all 397

public stations regardless of individual ownership, a simplification discussed further in Section 8.3. To support the daily planning process, the city area is additionally structured into 178 spatial zones using K-Means clustering. These zones provide a geographical abstraction that helps select promising daily work areas and organize station visits efficiently.




Figure 6: Three-phase benchmark for the number of zones k : mean CFA total cost by zone count. Each point represents the average over 100 simulation runs; colours indicate the benchmark phase.

The number of zones was determined through a three-phase benchmark rather than set arbitrarily. Phase 1 screened $k \in \{20, 40, \dots, 200\}$, Phase 2 refined the search around the promising range $k \in \{160, 170, \dots, 200\}$, and Phase 3 performed a fine-grained sweep over $k \in \{175, \dots, 185\}$. Each candidate was evaluated by the mean total simulation cost under CFA with value-based zone selection over 100 runs. Figure 6 shows the CFA cost curve across all three phases. Costs decrease steeply for small k and continue to fall until approximately $k = 178$ (33,919 €), beyond which they rise again slightly, indicating that finer zone granularity fragments the network without routing benefit. The value $k = 178$ was therefore selected as the operating parameter for all experiments.

Disruptions are modeled stochastically on a station-by-station and hour-by-hour basis. The probability of failure increases with time since the last maintenance intervention and is further modified by station-specific characteristics such as charger type and age. In this way, the simulation reflects both operational deterioration and heterogeneity across the charging network. The parameters governing this failure model, namely baseline rates, recovery curve shape, and station-level scaling factors, are not derived from historical maintenance records, as no consolidated failure history exists across Würzburg’s 83 station operators and no such data were available even from Stadtwerke Würzburg AG (WVW), the largest single operator. They were instead set to produce plausible failure frequencies and are therefore

synthetic. The implications of this for the interpretability of the results are discussed in Section 8.2.

7 Results and Analysis

This chapter evaluates the performance of the four policies developed in Chapter 5 across three experimental dimensions. Section 7.1 compares the four policy mechanisms under a controlled zone-selection setting to isolate the effect of the routing and prioritization logic. Section 7.2 then examines the role of zone selection itself, which turns out to be the dominant cost driver. Section 7.3 evaluates the VFA rollout layer, and Section 7.4 provides statistical validation of all key comparisons. All results are based on 500 Monte Carlo runs per variant using seeds 1–500. Because each seed generates an identical stochastic disturbance sequence across all variants, pairwise comparisons are matched by seed, which substantially reduces noise.

7.1 Policy Mechanism Comparison

Setup

Four policies are compared, namely Myopic (greedy cheapest insertion, with routing and drop decisions based on distance only and no value weighting), Myopic+ (routine tasks scored by power/distance and drop decisions by power/detour), CFA (routine tasks scored by $\tilde{C}$ /distance with learned θ and drop decisions by $\tilde{C} - \text{wage} \times \text{detour}$), and DB (same as CFA but with a state-dependent parameter δ that scales the distance exponent and the detour penalty). To isolate the effect of the routing and prioritization logic from zone-selection effects, all four policies are evaluated under the same centrality-based zone selection in this section. Centrality ranks zones by their geographic centrality within the network, measured by the mean distance from each zone’s centroid to all other zone centroids (Section 5.1), without any policy-specific value signal.

Results

Table 1 summarizes the results. Under centrality-based zone selection, all four policies perform at a similar level. Mean total costs range from 40,144 € (Myopic+) to 40,410 € (Myopic), a spread of just 266 € or 0.7 %. Myopic sits slightly above the three advanced policies, while Myopic+, CFA, and DB are virtually indistinguishable from one another (40,144–40,187 €, a difference of 43 €). Looking at the cost components, the advanced policies do reduce downtime costs relative to Myopic (16,642–16,702 € vs. 17,141 €, a reduction of 440–499 €), but this is offset by slightly higher labor costs (+169–188 €) and travel costs (+43–50 €), since power-weighted routing deviates from the nearest-neighbor order and thereby extends shift durations and travel distances marginally. The opposing shifts nearly cancel, leaving total costs within 0.7 % across all four policies. Section 7.4 confirms that none of these differences reach practical significance.

Table 1: Policy comparison under centrality-based zone selection (500 runs each). SD refers to the standard deviation of total costs across runs.

Policy	Mean Total (€)	SD (€)	Mean Labor (€)	Mean Travel (€)	Mean Downtime (€)	Mean Disrupt.
Myopic	40,410	5,134	22,098	1,171	17,141	162.3
Myopic+	40,144	5,197	22,286	1,216	16,642	163.2
CFA	40,187	5,156	22,271	1,221	16,695	162.0
DB	40,183	5,128	22,267	1,214	16,702	162.2

Discussion

The absence of meaningful policy differences under centrality is the central finding of this section. All four policies share the same greedy cheapest-insertion backbone and differ in how routine tasks are scored in the initial daily plan and how drop candidates are ranked during replanning. Under centrality-based zone selection, both teams are assigned zones by geographic centrality regardless of policy, so the station pool available to each team is essentially policy-independent. In this setting, a more sophisticated value-weighted scoring function yields no meaningful benefit, since the zone already determines which stations are available, and although the advanced policies do reduce downtime by visiting high-power stations earlier, this gain is offset by slightly higher labor and travel costs from deviating from a pure distance-minimizing order. The decisive question is therefore not *how* a policy routes within a zone, but *which* zone it selects.

Even Myopic, which assigns no priority to high-power stations in either routing or replanning, performs within 0.7 % of the advanced policies when zone selection is held constant. This directly motivates the analysis in Section 7.2.

7.2 Impact of Zone Selection on Cost

Setup

Section 7.1 showed that all four policies perform nearly identically under centrality-based zone selection. This section asks whether the choice of zone-selection strategy itself matters, and if so, by how much. Both zone-selection modes are evaluated for all four policies, namely centrality (zones ranked by geographic centrality within the network) and value-based (zones ranked by expected economic impact, computed via the policy's own value signal). All other simulation parameters are held constant.

Results

Table 2 reports mean total costs and cost components for all eight variants. Switching from centrality to value-based zone selection reduces mean total costs substantially across all four policies, by 5,525 € (13.7 %) for Myopic, 5,831 € (14.5 %) for Myopic+, 6,394 € (15.9 %) for CFA, and 6,512 € (16.2 %) for DB. In absolute terms, all four policies fall in the range

7 Results and Analysis

33,671–34,885 € under value-based selection, compared to 40,144–40,410 € under centrality.

Table 2: Cost comparison across zone-selection modes for all four policies (500 runs each). Δ is computed relative to centrality for each policy.

Policy	Zone Selection	Mean Total (€)	SD (€)	Mean Labor (€)	Mean Travel (€)	Mean Downtime (€)	Mean Disrupt.	Δ (%)
Myopic	centrality	40,410	5,134	22,098	1,171	17,141	162.3	–
Myopic	value-based	34,885	4,824	21,577	1,194	12,114	148.3	–13.7 %
Myopic+	centrality	40,144	5,197	22,286	1,216	16,642	163.2	–
Myopic+	value-based	34,313	4,727	21,546	1,211	11,556	146.9	–14.5 %
CFA	centrality	40,187	5,156	22,271	1,221	16,695	162.0	–
CFA	value-based	33,793	4,525	21,520	1,161	11,111	149.3	–15.9 %
DB	centrality	40,183	5,128	22,267	1,214	16,702	162.2	–
DB	value-based	33,671	4,546	21,491	1,158	11,022	148.9	–16.2 %

Figure 7 illustrates this directly, showing the average service day per zone for CFA under centrality-based selection (left) and value-based selection (right), both averaged over 500 runs. Under centrality, the service order is almost perfectly concentric, with centrally located zones visited early (green) and peripheral zones persistently late (red), regardless of their economic importance. Switching only the zone-selection mode breaks this pattern, since economically critical zones are pulled forward regardless of depot distance, and the concentric structure disappears.

Figure 8 shows the full cost distributions. The downward shift from centrality to value-based is consistent across all 500 runs and all four policies, confirming that the effect is not driven by a small subset of favorable scenarios. Labor costs fall only marginally under value-based selection (from approximately 22,200 € to 21,530 €, a reduction of approximately 3 %), driven by the 13–16 fewer disruptions per year: each disruption requires an unplanned insertion that extends the workday, so fewer disruptions translate directly into slightly shorter shifts and lower labor costs. Travel costs remain essentially constant at 1,160–1,220 € in both modes, since the total number of stations visited per day does not change, only their order. The entire reduction comes from downtime costs, which fall by 29–34 % across all policies. The disruption counts in Table 2 show that this is not only a matter of cheaper disruptions, since the absolute number of faults drops from roughly 162–163 under centrality to 147–149 under value-based selection, a reduction of approximately 13–16 disruptions per simulation year per policy. By pulling economically important zones forward in the service cycle, value-based selection keeps the days-since-maintenance of high-power stations low, which reduces their fault probability and thereby the total number of faults.

Discussion

The mechanism behind the zone-selection effect is straightforward, since value-based selection prioritizes zones containing stations with high rated power and long elapsed time




Figure 7: Mean service day per zone for CFA under centrality-based (left) and value-based (right) zone selection, averaged over 500 runs. Green = serviced early, red = serviced late.

Under centrality, the service order is strictly concentric, with centrally located zones always visited first and peripheral zones last. Value-based selection breaks this pattern, pulling economically critical zones forward to early service days regardless of their position in the network.

since last maintenance, which are precisely the stations generating the highest downtime costs per hour of unavailability. By visiting these zones earlier in the simulation, the policies resolve high-cost faults before they accumulate multiple days of downtime charges. The routing effort, labor and travel, is barely affected because the total number of stations visited per day remains the same. The disruption reduction (13–16 faults per simulation year) is a direct consequence of this mechanism, since stations that are maintained early see reduced days-since-maintenance, which lowers their per-hour failure probability via the stochastic failure model. Fewer failures accumulate because the maintenance cascade pre-empts them, not merely because each failure is handled faster.

A striking aspect of the results is that Myopic captures most of the zone-selection gain despite having no learned value function, saving 13.7 % versus 16.2 % for DB. This confirms that the dominant effect operates at the zone-selection level, not within the routing logic. Even a policy that is agnostic to station value at the routing stage captures most of the available gain simply by being directed toward the right zones. A residual gap remains under value-based selection, following a consistent hierarchy from Myopic (34,885 €) through Myopic+ (34,313 €) to CFA (33,793 €) and DB (33,671 €). CFA and DB outperform Myopic+ because their learned value signal selects qualitatively better zones, while Myopic+ relies only on a simpler power-recency product. The difference between CFA and DB is marginal,




Figure 8: Total cost distributions for all four policies under centrality (light) and value-based (dark) zone selection (500 runs each). Boxes show interquartile range, and whiskers extend to $1.5 \times \text{IQR}$.

since the additional distance-exponent parameter in DB only affects the late simulation phase when few stations remain unserved.

The residual gap has a concrete causal explanation. Although all four policies use value-based zone selection, they do so with structurally different value signals. Myopic scores zones by $\text{power} \times \text{recovery_curve}(dsm)$, a direct product of station power and maintenance recency. CFA and DB score zones via the learned approximation $\tilde{C}(k) = \theta^\top \phi_{\text{scaled}}(k)$, where ϕ incorporates four features, namely rated power, station age, recovery curve value, and mean distance to other stations. These different signals produce genuinely different zone assignments, not merely different within-zone orderings, but different geographic areas being prioritised on any given day. Within an assigned zone, the routing order compounds the effect, since CFA visits the highest- $\tilde{C}$ station first, while Myopic visits the nearest. The combined result is a *DSM cascade*, in which CFA reduces the days-since-maintenance of the highest-power stations more rapidly, lowering their failure probability in subsequent days. Because fault arrivals depend on dsm via the stochastic failure model, CFA teams face fewer high-cost disruptions in later days, not because they handle disruptions better, but because those disruptions occur less often.

Figure 9 illustrates this mechanism in a single run (seed 497). Both policies incur similar costs on days 1–2, when all stations start at maximum DSM and failure probabilities are uniformly high. From day 3, the CFA cumulative cost curve flattens relative to Myopic, reflecting a lower per-day fault rate on the high-power stations CFA has already visited.

7 Results and Analysis




Figure 9: Cumulative costs (top) and daily disruption counts (bottom) for Myopic and CFA under value-based zone selection, seed 497 (single run). The dashed vertical line marks day 9, where the gap widens most sharply. Myopic total: 52,491 €; CFA total: 37,191 €.

The gap becomes most visible on day 9, where Myopic encounters 9 disruptions and incurs 3,479 € in daily costs, while CFA sees only 4 disruptions and 1,275 €. This single day accounts for a large share of the eventual total gap of 15,300 € (Myopic: 52,491 €; CFA: 37,191 €). The disruption bar chart below the cumulative curves confirms that the advantage is not concentrated in one exceptional event, but sustained across the entire simulation, since after day 9, Myopic consistently faces more disruptions per day than CFA, and the cumulative cost lines never converge.

Compared to the policy differences observed in Section 7.1 (less than 1 % among all four policies under centrality), the zone-selection effect of 14–16 % is dramatically larger. Zone selection completely dominates the policy mechanism as a cost driver in this setting. Section 7.3 examines whether the VFA rollout layer can provide additional gains on top of value-based zone selection.

7.3 VFA Rollout: Performance and Limitations

Setup

Sections 7.1 and 7.2 established that value-based zone selection is the dominant cost driver, and that all four policies perform nearly identically under centrality-based zone selection. This section evaluates the VFA rollout layer, which wraps each base policy and overrides its drop decision whenever a short-horizon Monte Carlo evaluation identifies a better candidate. The rollout is evaluated on top of value-based zone selection for all four policies. Each rollout run uses the same seed as its corresponding base run, enabling paired cost comparisons. The rollout configuration follows Section 5.6: $k = 3$ candidates, 40-day horizon, 25 scenarios per candidate, 70 % win-rate threshold.

Table 3: Rollout effect per policy. All variants use value-based zone selection (500 runs each). Override rate = mean replan overrides / mean total carryover (≈ 48).

Policy	Without Rollout (€)	With Rollout (€)	Δ (€)	Δ (%)	Mean Overrides	Override Rate
Myopic	34,885	34,890	+5	+0.01 %	0.78	1.6 %
Myopic+	34,313	34,294	−19	−0.06 %	0.81	1.7 %
CFA	33,793	33,729	−64	−0.19 %	1.27	2.7 %
DB	33,671	33,627	−44	−0.13 %	1.20	2.5 %

Results

Table 3 summarizes the rollout effect for each policy. The rollout provides no meaningful benefit for any of the four policies. Myopic’s mean cost changes by only +5 € (+0.01 %), negligibly worse with the rollout. For the three advanced policies the effect is equally small, with Myopic+ saving 19 € (0.06 %), CFA saving 64 € (0.19 %), and DB saving 44 € (0.13 %). None of these differences exceed 0.2 % of mean total costs.

Figure 10 shows the paired cost difference (base minus rollout) per run for each policy. All four distributions are centred near zero with roughly symmetric spread, confirming that the rollout produces no systematic directional effect. The mean effects in Table 3 conceal substantial run-level variance, however. Across all four policies, approximately 22–24 % of individual runs produce strictly higher total costs with the rollout than without, with losses reaching up to 6,260 € in the most adverse case, while corresponding gains of comparable magnitude occur in 27–46 % of runs. The standard deviation of the paired cost difference ranges from ± 786 € (Myopic+) to ± 891 € (DB), roughly an order of magnitude larger than any mean effect.

Figure 11 shows the distribution of replan overrides per run. Across all four policies, most runs involve zero to two overrides, and runs with more than four overrides are rare. CFA and DB trigger more overrides on average (1.20–1.27) than Myopic and Myopic+ (0.78–0.81), since their value-weighted drop scores produce more differentiated candidate rankings that give the Monte Carlo evaluation more opportunities to disagree. Yet even the

7 Results and Analysis




Figure 10: Distribution of paired cost differences (base minus rollout) per run for each policy, shown as histograms with 250 € bins. Positive values indicate the rollout lowered cost, and negative values indicate the rollout increased cost. The dashed vertical line marks zero. Runs with $\Delta = 0$ (no override was triggered) are excluded from the histogram and reported separately in the annotation. The x-axis is clipped to $\pm 4,000$ €, and a small number of extreme values fall outside the plotted range. All variants use value-based zone selection (500 runs each).

more frequently overriding policies show no meaningful cost improvement.

Discussion

The absence of any rollout benefit has a clear structural explanation. The rollout intervenes exclusively at the *drop decision*. When a new disruption cannot be inserted without dropping a routine task, the rollout evaluates the top- k drop candidates via Monte Carlo lookahead and may override the base policy's choice. It does not affect how the initial daily plan is constructed.

For the advanced policies (Myopic+, CFA, DB), the Monte Carlo evaluation confirms rather than corrects the base policy's drop choice, since their drop-score functions already incorporate station value, so even when the rollout disagrees (averaging 0.8–1.3 overrides per run out of ≈ 48 carryover decisions) the cost consequence is too small to accumulate into a measurable aggregate difference.

7 Results and Analysis




Figure 11: Distribution of replan overrides per run for all four rollout variants (500 runs each). Dashed vertical line shows the mean.

Myopic’s situation is different in structure but leads to the same outcome. Myopic’s drop-score uses $1/\text{detour}$, routing efficiency only, ignoring station power, and one might expect the rollout to systematically correct this. In practice, however, only 0.78 overrides per run occur, and they yield no net improvement. A contributing factor is that the rollout cannot address Myopic’s initial-plan scoring function, which also ignores station power ($1/\text{distance}$ versus Myopic+’s power/distance). Even if every drop decision were perfectly corrected, the ordering of routine stations within the initial plan would remain power-agnostic. The small remaining gap between Myopic and the advanced policies, as will be confirmed in Section 7.4 ($d = 0.17\text{--}0.37$), likely reflects this combined effect.

The run-level variance visible in Figure 10 warrants a closer look at why the rollout sometimes produces substantially worse outcomes. In approximately 22–24 % of runs across all policies, the rollout increases total costs, occasionally by several thousand euros. The mechanism is finite-sample estimation noise. The rollout evaluates each drop candidate over 25 Monte Carlo scenarios spanning a 40-day horizon. A measured Win-Rate of, say, 76 % for a given candidate means that 19 of 25 sampled futures favoured that choice. The actual simulation, however, is a single draw from all possible futures, and if that draw falls in the 24 % minority, the override is wrong. With only $n = 25$ scenarios, a coin-flip candidate would produce a measured Win-Rate of 76 % or above in roughly 0.7 % of evaluations by chance alone, whereas a Win-Rate of 64 % or above already occurs in about 11 % of evaluations, illustrating how easily moderate-looking win rates can arise by chance at this sample size. This structural noise cannot be eliminated without sharply increasing the number of scenarios.

A concrete run illustrates the cascading consequence of a single erroneous override. In seed 6 of the Myopic rollout, exactly one override occurs over the entire simulation. The mechanism unfolds as follows. On day 11 at 9:00, a minor Type-1 fault is reported. The base policy drops a routine stop from one team’s route and inserts the fault there. The rollout

evaluates the top drop candidates in 25 Monte Carlo scenarios and finds that dropping a different stop, located in the other team's route, appears superior in 76 % of the sampled futures—an apparently strong signal. It overrides the base choice, so the fault is inserted into the other team's route instead, shifting that team's remaining stops back by roughly one hour.

Both runs then face the identical disruption sequence for the remainder of the day (same seed), but from different route states. At 14:00, a Type-1 fault at a 250 kW fast-charging station is reported. The base run still accommodates it by dropping two further routine stops, whereas in the rollout run the shifted schedules admit no feasible insertion, so the high-power fault carries over to the next day. The backlog then compounds: on day 12, the base run resolves four of its seven incoming disruptions, while the rollout run resolves only two and accumulates five carryovers. Costs diverge sharply, with the rollout run incurring 907 € additional cost on day 12 and 983 € on day 13, and the accumulated backlog eventually extends the rollout run by a full working day (43 instead of 42 days to complete all maintenance). The total excess cost over the full simulation is 5,139 € (45,572 € vs. 40,433 €), generated entirely by a single decision that appeared correct in 76 % of the sampled scenarios. Table 4 summarizes this sequence.

Table 4: Seed 6 cascade: cumulative total costs around the single day-11 override (Myopic with value-based zone selection, base vs. rollout run).

Day	Event	Base (€)	Rollout (€)	Excess (€)
1–10	Identical decisions and costs	15,633	15,633	0
11	Override at 9:00 (win rate 76 %); 250 kW fault becomes infeasible at 14:00 and carries over	17,435	17,511	+75
12	Base resolves 4 of 7 disruptions, rollout only 2 of 7	18,977	19,959	+982
13	Rollout works off accumulated backlog	20,382	22,347	+1,965
End	Full simulation (42 vs. 43 days)	40,433	45,572	+5,139

Seed 6 is a particularly transparent case because exactly one override occurs, making the causal chain traceable. In runs with multiple overrides, the effects of individual decisions intermix and partially cancel, producing the broad, symmetric distribution visible in Figure 10 rather than a single sharp divergence.

Taken together, the rollout behaves as a high-variance, zero-expectation intervention. For Myopic and Myopic+, roughly 47–49 % of runs see no change at all because no override is triggered, whereas for CFA and DB this share is lower (30–33 %) due to their higher override rates. Across all policies, around 22–24 % of runs incur higher costs, occasionally by several thousand euros, while 27–46 % benefit slightly. The mean improvement is negligible in all

cases, but the downside risk is substantial. This asymmetry, combined with the computational overhead of Monte Carlo evaluation at every disruption event, means the rollout does not justify its cost when paired with value-based zone selection. The dominant lever identified in Section 7.2 essentially saturates the available improvement potential before the rollout can contribute. Statistical validation of all findings in Sections 7.1–7.3 is provided in Section 7.4.

7.4 Statistical Validation

Setup

All comparisons in Sections 7.1–7.3 are validated using paired statistical tests. Because each seed generates an identical stochastic disturbance sequence across all variants, runs sharing the same seed form natural pairs. This pairing eliminates the between-run variance due to the stochastic environment and substantially increases test power. Two complementary tests are applied, namely a paired t -test (parametric, assumes normally distributed differences) and a Wilcoxon signed-rank test (non-parametric, robust to non-normality and outliers). With $n = 500$ paired observations, even very small differences can yield statistically significant p -values. Cohen's d (standardized mean difference of paired deltas) is therefore used as the primary criterion for practical relevance: $d < 0.20$ is considered negligible, 0.20 – 0.50 small, 0.50 – 0.80 medium, and $d > 0.80$ large.

Results

Table 5 reports the key paired comparisons. Four distinct patterns emerge.

Zone selection ($d = 1.29$ – 1.50): The strongest effect in the entire study. For all four policies, switching from centrality to value-based zone selection produces differences of 1.3 – 1.5 standard deviations, significant at $p < 0.001$ under both tests. This is the only effect with a large Cohen's d .

Policy differences under centrality ($d = 0.00$ – 0.07): All comparisons are statistically non-significant ($p > 0.10$) with negligible effect sizes. No policy outperforms any other in a practically meaningful way when zone selection is held constant at centrality.

Policy differences under value-based zone selection ($d = 0.05$ – 0.37): Myopic differs significantly from all three advanced policies ($d = 0.17$ – 0.37 , $p < 0.001$). Myopic+ also differs significantly from CFA and DB ($d = 0.16$ – 0.20 , $p \leq 0.002$ under both tests), though the effect sizes are negligible by Cohen's convention. The one non-significant comparison is CFA vs. DB ($d = 0.05$, $p = 0.26$), confirming these two policies are effectively equivalent once zone selection is optimized. The DB parameter δ was designed to increase routing compactness in the late simulation phase (once fewer than 130 stations remain unserved), with the intention of completing routes faster and thereby reducing labor costs. The data confirm this mechanism is active but marginal, as DB saves 29 € in mean labor costs relative to CFA (21,491 € vs. 21,520 €), a difference of 0.1 %. The effect is limited because the δ -

Table 5: Paired statistical tests for key comparisons (500 matched seeds each). Cohen’s d is the primary criterion for practical relevance, and p -values are from the paired t -test and Wilcoxon signed-rank test.

Comparison	Mean Δ (€)	d	t -test p	Wilcoxon p	Finding
<i>Zone selection: centrality $\rightarrow$ value-based</i>					
Myopic	−5,525	1.29	< 0.001	< 0.001	*** large
Myopic+	−5,831	1.37	< 0.001	< 0.001	*** large
CFA	−6,394	1.50	< 0.001	< 0.001	*** large
DB	−6,512	1.49	< 0.001	< 0.001	*** large
<i>Policy comparison under centrality</i>					
Myopic vs. Myopic+	+266	0.07	0.128	0.137	n.s.
Myopic vs. CFA	+223	0.06	0.181	0.157	n.s.
Myopic vs. DB	+227	0.06	0.218	0.369	n.s.
CFA vs. DB	+4	0.00	0.980	0.732	n.s.
<i>Policy comparison under value-based zone selection</i>					
Myopic vs. DB	+1,214	0.37	< 0.001	< 0.001	** small
Myopic vs. CFA	+1,092	0.31	< 0.001	< 0.001	** small
Myopic vs. Myopic+	+573	0.17	< 0.001	< 0.001	* negligible
Myopic+ vs. DB	+642	0.20	< 0.001	< 0.001	* negligible
Myopic+ vs. CFA	+520	0.16	< 0.001	0.002	* negligible
CFA vs. DB	+122	0.05	0.258	0.568	n.s.
<i>Rollout effect (value-based base $\rightarrow$ +Rollout)</i>					
Myopic	+5	0.01	0.900	0.310	n.s.
Myopic+	−19	0.02	0.590	0.626	n.s.
CFA	−64	0.08	0.073	< 0.001	negligible [†]
DB	−44	0.05	0.268	< 0.001	negligible [†]

*** large ($d > 0.80$); ** small ($0.20 \leq d < 0.50$); * statistically significant but negligible ($d < 0.20$); n.s. = not statistically significant ($p > 0.05$). [†] Wilcoxon detects a statistically significant but practically negligible systematic shift ($d < 0.10$).

switch affects only the final phase of the simulation, and because shift duration is primarily determined by the number of stations and disruptions rather than routing compactness.

Rollout ($d = 0.01$ – 0.08): No rollout comparison exceeds $d = 0.10$. For Myopic and Myopic+ the t -test and Wilcoxon both return non-significant results. For CFA and DB, the Wilcoxon detects a statistically significant but practically negligible systematic shift ($d = 0.08$ and $d = 0.05$ respectively), a consequence of the large sample size ($n = 500$) rather than an economically meaningful effect.

Figure 12 summarizes all comparisons by Cohen’s d . The zone-selection effects stand apart clearly, while all other effects fall below the $d = 0.50$ threshold, and most are below $d =$

7 Results and Analysis

0.20.




Figure 12: Cohen's d for all key comparisons. Dashed vertical lines mark conventional thresholds (small: 0.20, medium: 0.50, large: 0.80). Zone-selection effects are the only comparisons with a large effect size, and all other comparisons fall below the medium threshold ($d = 0.50$).

Discussion

The statistical analysis confirms and sharpens the findings from Sections 7.1–7.3. Three conclusions stand out.

First, zone selection is the only intervention with a large and robust effect. The $d > 1.2$ values for all four policies place this effect firmly in the “large” category by any standard, and the consistency across policies (all four show similar magnitudes) rules out policy-specific confounders. With $n = 500$ matched pairs, the p -values are effectively zero, and both parametric and non-parametric tests agree.

Second, under value-based zone selection, policy differences become detectable though remain small. Myopic differs significantly from all three advanced policies ($d = 0.17$ – 0.37), and Myopic+ differs from CFA and DB ($d = 0.16$ – 0.20), even though all four policies were statistically indistinguishable under centrality. The residual gap has two interacting

sources, both explained in Section 7.2. On the one hand, even under value-based mode, the policies rely on structurally different value signals, so they select genuinely different zones. Myopic uses a simple power-recency product, while CFA and DB use the learned four-feature approximation $\tilde{C}$. On the other hand, within an assigned zone, Myopic’s power-agnostic scoring ($1/\text{distance}$, $1/\text{detour}$) visits and protects stations in a different order than CFA’s value-weighted scoring. Together these differences drive a DSM cascade, since CFA reduces days-since-maintenance on the highest-power stations faster, pre-empting the costliest faults. Once zone selection directs teams toward high-value areas, this compound mechanism becomes visible, but the resulting effect sizes ($d = 0.16\text{--}0.37$) are small to negligible in practical terms.

Third, the absence of a meaningful rollout effect is confirmed statistically, not just descriptively. For Myopic and Myopic+, neither test detects any systematic improvement. The Wilcoxon does flag CFA and DB rollouts as statistically significant, but with $d < 0.10$, this reflects the sensitivity of large-sample tests to trivially small effects, not a practically relevant finding. Given the computational cost of Monte Carlo rollouts at every disruption event, the rollout layer does not justify its overhead in this setting.

Taken together, DB with value-based zone selection and rollout achieves the lowest mean total cost (33,627 €) across all tested configurations and therefore represents the best-performing setup overall. The underlying gap between DB and CFA under value-based zone selection is negligible ($d = 0.05$, $p = 0.26$, Table 5), and the rollout adds no meaningful further separation. The decisive factor is zone selection: DB with centrality-based zone selection is outperformed by Myopic with value-based zone selection by a wide margin.

8 Discussion

This chapter contextualizes the empirical findings from Chapter 7 within the theoretical framework and operational constraints. The discussion proceeds in three parts: interpretation of the results and their policy implications, analysis of the model’s limitations and assumptions, and assessment of practical applicability to other networks.

8.1 Interpretation of Results

The results reveal a clear hierarchy of effects. Zone selection is the dominant cost driver ($d = 1.29\text{--}1.50$, -14 to -16%), dwarfing both policy differences ($d \leq 0.37$) and the VFA rollout ($d \leq 0.08$). Under centrality-based zone selection, all four policies are statistically indistinguishable, whereas once zone selection is optimized, Myopic falls behind the advanced policies by a small but detectable margin ($d = 0.17\text{--}0.37$), while CFA and DB remain effectively equivalent ($d = 0.05$). The VFA rollout produces no economically meaningful improvement for any policy. The practical implication is unambiguous, namely that a well-chosen zone assignment with the simplest policy (Myopic) outperforms a poorly chosen zone assignment with the most sophisticated policy (DB) by a wide margin. Investment in

zone-selection quality yields far greater returns than refinement of the routing heuristic.

8.2 Limitations

The results reported in this thesis rest on a set of modelling and data assumptions that bound both their validity and their generalizability. This section identifies the most consequential limitations.

The entire experimental design assumes one central depot in Würzburg and 397 spatially concentrated charging stations. All routing and clustering logic is calibrated to this geometry. Operators with multiple depots, distributed vehicle fleets, or spatially sparse networks would face a substantially different problem structure. The K-Means zone partitioning and the greedy cheapest-insertion backbone are not inherently limited to single-depot settings, but the centrality-based zone scoring, which favours zones central to a single contiguous service area, and the depot-anchored route construction would require rethinking in a multi-depot context.

The stochastic failure simulation uses a parametric model in which failure probability is a function of days since last maintenance and rated station power, governed by four scalar parameters (p_1 , p_2 , λ , and the initial recovery factor). These parameters were set to produce plausible failure rates rather than estimated from historical records. Würzburg's 397 public stations are operated by 83 distinct providers, and no consolidated failure history exists across the network. An attempt to obtain operational data for calibration was made with Stadtwerke Würzburg AG (WVV), the largest single operator with 36.5 % of stations, but did not receive a response. The quantitative cost figures reported in Chapter 7 therefore reflect the synthetic distribution rather than actual field conditions. Relative comparisons between policies remain valid under this assumption, since all policies face identical failure sequences, but absolute cost magnitudes should not be read as predictions of real operating costs.

The model fixes the number of maintenance teams at two and the daily service capacity at 20 stations per team. In practice, operators may adjust staffing levels, deploy temporary support teams during peak fault periods, or subcontract overflow tasks. None of these options are represented in the current model. Service times are likewise modeled as homogeneous per fault type, whereas actual durations vary substantially with fault severity and technician skill. Similarly, the model does not account for vehicle range constraints, spare part availability, or technician skill heterogeneity, all of which affect real-world maintenance planning.

Hourly traffic matrices are used to represent time-dependent travel conditions, computed from the OSRM routing engine using OpenStreetMap data. While this is a reasonable approximation, OSRM does not model real-time congestion, road closures, or seasonal variation. The routing component therefore represents a simplified but not unrealistic version of actual travel conditions in Würzburg.

The rollout layer evaluates the top- k drop candidates via Monte Carlo simulation at each

disruption event. In the current setting ($k = 3$, 40-day horizon, 25 scenarios per candidate), this adds modest computational overhead. However, the cost scales linearly with the number of disruptions per day, the number of candidates k , and the simulation depth. For operators with larger fleets, higher disruption rates, or tighter replanning time budgets, the overhead could become prohibitive. The negligible performance benefit observed here makes this a moot point for the current problem instance, but scalability remains a concern if the approach is extended to more dynamic settings.

The optimization target is a scalar cost composed of labor, travel, and downtime cost components. In practice, maintenance operators face multi-objective trade-offs, since service level agreements may impose hard response-time constraints, equity considerations may require balanced workloads across teams, and customer satisfaction may be a metric distinct from pure economic downtime cost. The model does not represent these objectives, and results may differ under alternative objective specifications.

8.3 Practical Relevance and Transferability

The thesis was motivated by a concrete operational problem, namely how a maintenance team for a growing EV charging network should plan and adapt its daily routes under uncertainty. Würzburg's public charging network, 397 stations distributed across 83 commercial operators, with the municipal utility Stadtwerke Würzburg AG (WV) as the single largest operator at 36.5 % of stations, served as the reference case. Since no single operator maintains the network as a whole, the model treats the two maintenance teams as a hypothetical, city-wide coordinating entity rather than as a description of WV's actual internal operations. Such a consolidated maintenance service is a plausible arrangement (for example, contracted out to a third-party provider) but does not currently exist in this form. An attempt to obtain operational data and benchmark figures from WV, as the largest individual operator, was made during the preparation of this thesis but received no response. A direct empirical comparison between the modelled policies and any operator's current scheduling practice is therefore not possible. The following discussion addresses the practical relevance of the findings at a conceptual level and considers the conditions under which the approach could be deployed or transferred to other settings.

Deploying the system in a real maintenance operation would require three types of input data that are not always readily available. Historical failure records per station are needed to calibrate the stochastic failure model, specifically to estimate per-station failure rates and the relationship between maintenance recency and subsequent failure probability. Station-level metadata (rated power, installation date, precise GPS coordinates) must be available in structured form. Finally, real-time fault notifications must be routable to the planning system, either via sensor telemetry or a fault-reporting interface. None of these requirements are technically prohibitive for an operator managing a comparably sized network of around 400 stations, but they imply a data management infrastructure that may not exist in its current form for any individual provider in the Würzburg market.

The approach is best suited to operators with a centralized depot structure, a stable fleet of

two to four maintenance teams, and a medium-density station network in a single metropolitan area. These conditions match a municipal utility such as Stadtwerke Würzburg AG (WVV) or a third-party service consolidating maintenance across several of Würzburg’s other operators, and also describe many municipal or regional charging network operators in Germany. The core computational cost is low, since the greedy routing policies run in milliseconds, and even the zone-selection logic adds negligible overhead. The only computationally intensive component, the VFA rollout, was shown to provide no meaningful benefit in this setting and can be omitted without loss.

Three structural features of the current model constrain transferability to other contexts. The zone-assignment logic assumes a single depot from which both teams depart, so operators with spatially distributed vehicle pools would need to replace the centrality-based zone ranking with a fleet-assignment-aware proximity measure. The CFA cost approximation $\hat{C}$ was trained on synthetic data from the Würzburg network, and its learned weights θ are unlikely to transfer directly to networks with substantially different station densities, power distributions, or failure dynamics. Retraining on local data is straightforward given the offline regression procedure, but requires sufficient historical observations. The model also assumes homogeneous service times, so operators with a mix of hardware generations or station types requiring different maintenance procedures would need to extend the time-window logic accordingly.

Despite these constraints, the central finding, that zone selection quality dominates routing policy sophistication, is likely to transfer beyond the specific experimental setting. Its dominance reflects a structural property of capacity-constrained routing under uncertainty: when teams have limited daily capacity, the choice of *which areas to serve* outweighs *how to route within them*. Since this argument does not depend on Würzburg’s specific geometry, power distribution, or failure rates, it should apply to any urban charging network with spatial clustering, variable power ratings, and limited team capacity. It implies that a maintenance operator’s most impactful intervention is not algorithmic refinement of the routing heuristic, but rather the development of a zone-value signal that incorporates station power and maintenance recency. This information is typically available to operators from internal maintenance records alone. A simple rule-based implementation of value-based zone selection, without any of the ADP machinery developed in this thesis, could plausibly capture a substantial fraction of the cost improvement observed here. The more sophisticated policies (CFA, DB) provide incremental benefits that are statistically detectable but economically modest, on the order of 1–2 % relative to a well-configured Myopic+ baseline under value-based zone selection.

These considerations suggest a phased deployment path. A first phase implements rule-based value-based zone selection using the power–recency product $p_i \cdot \rho(dsm_i)$, capturing the dominant share of the observed savings (13–16 %) with low implementation risk and no learned components. A second phase upgrades the routing policy from Myopic to Myopic+ or CFA, the latter requiring retraining of θ on local historical data, and adds the small incremental gains quantified in Section 7.4. A third phase, the VFA rollout, is not recom-

mended, since it adds negligible expected improvement at meaningful computational cost and with substantial run-level downside risk (Section 7.3).

9 Conclusion and Outlook

This final chapter concludes the thesis. It summarizes the key findings (Section 9.1), articulates the methodological and practical contributions (Section 9.2), identifies open research questions and directions for future work (Section 9.3), and closes with concluding remarks. The limitations of the work and the transferability of its findings were discussed in Chapter 8.

9.1 Summary of Key Findings

This thesis addressed the research question of how maintenance operations for electric vehicle charging stations can be planned dynamically under uncertainty, so as to balance downtime, travel costs, and resource utilization efficiently. The answer, developed through a route-based MDP and a hierarchy of ADP policies, reveals a clear and surprising structure: *the choice of which geographic zones to prioritize dominates all other decision factors by a dramatic margin.*

The empirical results from 500 Monte Carlo runs per configuration expose a strict hierarchy of cost drivers. Zone selection is the overwhelming leverage point. When switching from centrality-based (geographic) to value-based (economic) zone selection, total costs drop by 5,525 € to 6,512 € (13.7 – 16.2 %) across all four policies, with Cohen’s $d = 1.29\text{--}1.50$ (large effect). This zone-selection impact dwarfs all policy improvements. By contrast, routing policy sophistication matters little. Under centrality-based zone selection, all four policies (Myopic, Myopic+, CFA, DB) perform within 0.7 % of each other (40,144–40,410 €), with $d < 0.07$ (negligible). The implication is clear. With bad zone selection, policy choice makes almost no difference.

However, once zone selection is optimized (value-based), policy differences become detectable. Myopic costs 34,885 €, while the advanced policies (CFA, DB) cost 33,671 €–33,793 €, a gap of 1,092 €–1,214 € (3.2 – 3.6 %, $d = 0.17\text{--}0.37$). These differences, while statistically significant with large sample sizes ($n = 500$), remain small in practical terms. The VFA rollout layer adds negligible value (0.01 – 0.19 % improvement, $d \leq 0.08$) and should be omitted.

The core practical message is unambiguous: *zone selection quality dominates policy sophistication.* The Myopic policy paired with value-based zone selection (34,885 €) outperforms DB paired with centrality-based zone selection (40,183 €) by 5,298 € annually. When zone selection is already optimized, the choice between policies becomes a trade-off: DB achieves 33,671 € (1,214 € better than Myopic under the same zone selection), but requires learned value-function parameters, while Myopic+ reaches 34,313 € (642 € more than DB) using only power-weighted scoring.

The mechanism underlying the zone-selection dominance is a “days-since-maintenance cascade”. Value-based zone selection prioritizes high-power stations with long elapsed maintenance intervals—precisely those generating the highest downtime costs. By serving these zones early, teams prevent fault accumulation. Fewer disruptions occur not because individual repairs are faster, but because the maintenance sequence pre-empts failures before they arise. This mechanism operates independently of within-zone routing order, explaining why all four policies capture nearly identical zone-selection gains despite different value functions.

9.2 Contributions of This Thesis

This thesis contributes to maintenance optimization, stochastic vehicle routing, and approximate dynamic programming through methodological and practical advances. Methodologically, the problem is formalized as a route-based Markov decision process where actions modify complete route plans rather than selecting individual next stops, capturing how maintenance actually unfolds in a dual-mode structure (initial planning and intraday re-planning) absent from classical Vehicle Routing Problem literature. The thesis develops a structured progression of policies from Myopic (distance-only baseline) through Myopic+ (power-weighted), Cost Function Approximation (learned value function via contrastive regression), Dynamic Balance (state-dependent routing parameters), and Value Function Approximation (Monte Carlo rollout), enabling systematic evaluation of whether sophisticated methods justify their added complexity. Zone selection itself is formalized as a daily structural decision distinct from within-zone routing, allowing its impact to be isolated and revealing that it dominates policy choice. The empirical validation rests on 500 seed-matched Monte Carlo runs per configuration with rigorous statistical testing (paired t -tests, Wilcoxon, Cohen’s d) providing evidence of effect sizes uncommon in operations research.

On the practical side, this thesis delivers concrete deployment guidance. Value-based zone selection, implementable rule-based without machine learning, delivers annual savings of 5,500–6,500 € (14–16 %) for Würzburg-scale networks. Policy refinement (upgrading from Myopic to CFA or DB) adds only a further 500–1,200 € and should be considered only if infrastructure complexity is acceptable. The VFA rollout should be skipped entirely: it provides no benefit and introduces high downside risk. A maintenance operator can implement the core mechanism (value-based zone selection) using a spreadsheet and station-level data (power, days-since-maintenance), making the primary finding immediately actionable without specialized software or machine learning expertise.

9.3 Open Research Questions and Future Directions

Future work could extend this thesis in several directions. Operationally, adaptive zone geometry could allow K-Means boundaries to shift based on current DSM distribution (re-clustering every 5–10 days rather than holding zones fixed), potentially improving upon the static partitioning used here. Sensor-based predictive maintenance represents a promising direction: equipping stations with sensors and using ML to predict failures days in advance

could enable proactive scheduling and reduce reactive disruptions, though it requires substantial hardware investment and validation. Real-time learning loops could retrain policy parameters (CFA weights θ , zone-selection weights) based on new operational data, requiring automated data pipelines and A/B testing infrastructure to validate improvements. Extension to multi-depot environments with dynamic team reassignment is feasible (zones and scoring generalize directly), though team-to-zone assignment becomes more complex; inter-depot coordination could reduce empty positioning trips.

Theoretically, several open questions remain. Is zone-selection dominance fundamental to all capacity-constrained routing problems under uncertainty, or specific to this problem's structure (small teams, high disruption rates)? What Monte Carlo sample size is actually needed for rollout to outperform greedy decision-making? And would different policies converge to the same strategy if each were free to learn optimal zone assignments, or do genuinely different solutions emerge? These questions would deepen understanding of the structure of service planning under uncertainty and guide future algorithm development.

9.4 Concluding Remarks

This thesis set out to develop sophisticated approximate dynamic programming policies for the dynamic maintenance routing problem of electric vehicle charging stations. The central hypothesis was that combining learned value functions, anticipatory rollout evaluation, and state-aware routing would yield substantial cost improvements. The empirical investigation revealed a surprising and more nuanced reality: the dominant cost driver is not policy sophistication but rather the choice of which geographic zones to prioritize for service.

This finding is not a failure of the ADP framework but rather a validation of the scientific method applied rigorously. The thesis began with a theoretically motivated hierarchy of policies, measured their performance carefully while controlling for stochastic noise through seed-matched experiments, and isolated individual effects through systematic design, first holding zone selection constant and then holding the policy constant. This process revealed that the initial focus on policy sophistication was misplaced: the real lever for cost reduction lies upstream, in zone selection. For future research in dynamic service planning, the lesson is clear: *empirical effect-size measurement must precede algorithm development*. A well-measured simple heuristic outperforms an un-validated sophisticated policy. This thesis provides both the methodological framework to measure and the empirical evidence to guide such prioritization.

For maintenance operators contemplating deployment, the practical message is equally clear. Investment in zone-selection infrastructure—reliable data on station power and maintenance recency, clear scoring rules based on economic impact, and daily zone-selection execution—delivers substantially greater returns than refinement of routing algorithms, and a phased deployment path toward this goal is outlined in Section 8.3. When choosing among competing technologies and initiatives, the operator should ask: “Does this improve zone selection quality?” If yes, pursue it. If it only improves within-zone routing, defer it in

favor of other priorities.

The work also contributes a methodological framework—the route-based MDP, the offline-online ADP structure, and the systematic policy hierarchy—that should serve future work in dynamic service planning. The framework is not specific to charging infrastructure but applies to field service, emergency response, and many other domains where teams must plan routes, adapt to disruptions, and balance cost and service quality under uncertainty. By providing both theory and evidence, the thesis aims to advance not only the specific problem of charging-station maintenance but also the broader field of dynamic service planning under uncertainty.

References

- Abbatecola, Lorenzo, Maria Pia Fanti, and Walter Ukovich (2016). "A review of new approaches for dynamic vehicle routing problem". In: *2016 IEEE International Conference on Automation Science and Engineering (CASE)*. IEEE, pp. 361–366.
- Akhmetbek, Yernar and Taukekhan Mustakhov (2023). "Stochastic dynamic vehicle routing problem survey". In: *SDU Bulletin: Natural and Technical Sciences* 62.2, pp. 75–87.
- Alaswad, Suzan and Yisha Xiang (2017). "A Review on Condition-Based Maintenance Optimization Models for Stochastically Deteriorating System". In: *Reliability Engineering & System Safety* 157, pp. 54–63.
- Anuar, Wadi Khalid, Lai Soon Lee, Hsin-Vonn Seow, and Stefan Pickl (2022). "A multi-depot dynamic vehicle routing problem with stochastic road capacity: An MDP model and dynamic policy for post-decision state rollout algorithm in reinforcement learning". In: *Mathematics* 10.15, p. 2699.
- Basso, Rafael, Balázs Kulcsár, Ivan Sanchez-Diaz, and Xiaobo Qu (2022). "Dynamic stochastic electric vehicle routing with safe reinforcement learning". In: *Transportation research part E: logistics and transportation review* 157, p. 102496.
- Bundesministerium der Finanzen (2024). *Reisekosten: Kilometerpauschale für Dienstreisen*. URL: <https://www.bundesfinanzministerium.de> (visited on 11/01/2024).
- Bundesverband der Energie- und Wasserwirtschaft (2024). *BDEW-Ladeinfrastruktur-Monitor*. URL: <https://www.bdew.de/energie/elektromobilitaet-daten-und-fakten/> (visited on 11/01/2024).
- De Jonge, Bram and Philip A. Scarf (2020). "A Review on Maintenance Optimization". In: *European Journal of Operational Research* 285.3, pp. 805–824.
- Dekker, Rommert (1996). "Applications of Maintenance Optimization Models: A Review and Analysis". In: *Reliability Engineering & System Safety* 51.3, pp. 229–240.
- Hildebrandt, Florentin D, Barrett W Thomas, and Marlin W Ulmer (2023). "Opportunities for reinforcement learning in stochastic dynamic vehicle routing". In: *Computers & operations research* 150, p. 106071.
- jobvector GmbH (2024). *Gehalt Elektroinstallateur Bayern*. URL: <https://www.jobvector.de/gehalt/elektroinstallateur/bayern/> (visited on 11/01/2024).
- Mardešić, Nikola, Tomislav Erdelić, Tonči Carić, and Marko Đurasević (2023). "Review of stochastic dynamic vehicle routing in the evolving urban logistics environment". In: *Mathematics* 12.1, p. 28.
- Molęda, Marek, Bożena Małysiak-Mrozek, Weiping Ding, Vaidy Sunderam, and Dariusz Mrozek (2023). "From corrective to predictive maintenance—A review of maintenance approaches for the power industry". In: *Sensors* 23.13, p. 5970.
- Pillac, Victor, Michel Gendreau, Christelle Guéret, and Andrés L Medaglia (2013). "A review of dynamic vehicle routing problems". In: *European Journal of Operational Research* 225.1, pp. 1–11.
- Ritzinger, Ulrike, Jakob Puchinger, and Richard F Hartl (2016). "A survey on dynamic and stochastic vehicle routing problems". In: *International Journal of Production Research* 54.1, pp. 215–231.

-
- Stadt Würzburg (2024). *Ladesäulen in Würzburg – Open Data Portal*. URL: <https://wuerzburg.opendatasoft.com/explore/dataset/10001wuerzburg-stadt/information/> (visited on 06/20/2026).
- Stein, Jonas, Florentin D Hildebrandt, Marlin W Ulmer, and Barrett W Thomas (2025). “Learning state-dependent policy parametrizations for dynamic technician routing with rework”. In: *Transportation science* 59.5, pp. 1153–1171.
- Ulmer, Marlin W (2017). *Approximate dynamic programming for dynamic vehicle routing*. Vol. 61. Springer.
- (2020). “Horizontal combinations of online and offline approximate dynamic programming for stochastic dynamic vehicle routing.” In: *Central European Journal of Operations Research* 28.1, p. 279.
- Ulmer, Marlin W, Justin C Goodson, Dirk C Mattfeld, and Marco Hennig (2019). “Offline–online approximate dynamic programming for dynamic vehicle routing with stochastic requests”. In: *Transportation Science* 53.1, pp. 185–202.
- Ulmer, Marlin W, Justin C Goodson, Dirk C Mattfeld, and Barrett W Thomas (2020). “On Modeling Stochastic Dynamic Vehicle Routing Problems”. In: *EURO Journal on Transportation and Logistics* 9.2, p. 100008. DOI: 10.1016/j.ejtl.2020.100008.
- Ulmer, Marlin W, Dirk C Mattfeld, and Felix Köster (2018). “Budgeting time for dynamic vehicle routing with stochastic customer requests”. In: *Transportation Science* 52.1, pp. 20–37.
- Yenigun, Onur, Gozde Karatas Baydogmus, and Kazim Yildiz (2025). “Smart Routing for EV Charge Point Operators in Mega Cities: Case Study of Istanbul”. In: *arXiv preprint arXiv:2509.21369*.
- Zhang, Meng, Lulu Sun, Qi Jiang, and Wei Cui (2024). “Routing Optimization for Large-Scale Inspection-Maintenance of EV Charging Spots With Multiple Types of Personnel and Composite Customer Satisfaction: A Multi-Stage Approach”. In: *IEEE Access* 12, pp. 158189–158206. DOI: 10.1109/ACCESS.2024.3486686.
- Zou, Han and Maged M Dessouky (2018). “A look-ahead partial routing framework for the stochastic and dynamic vehicle routing problem”. In: *Journal on Vehicle Routing Algorithms* 1.2, pp. 73–88.

List of Acronyms

ADP	Approximate Dynamic Programming
CBM	Condition-Based Maintenance
CFA	Cost Function Approximation
DB	Dynamic Balance
DP	Dynamic Programming
DSM	Days Since Maintenance
EV	Electric Vehicle
MDP	Markov Decision Process
OSRM	Open Source Routing Machine
RL	Reinforcement Learning
SD	Standard Deviation
SDVRP	Stochastic Dynamic Vehicle Routing Problem
TSP	Travelling Salesman Problem
VFA	Value Function Approximation
VRP	Vehicle Routing Problem
WVV	Würzburger Versorgungs- und Verkehrs-GmbH

A Appendix A

Table 6: Summary of notation used in Chapter 4

Symbol	Description
<i>Sets and indices</i>	
$I = \{1, \dots, 397\}$	Set of charging stations
$K = \{1, 2\}$	Set of service teams
$V = \{0\} \cup I$	Complete node set (depot + stations)
$Z = \{1, \dots, N_Z\}, N_Z = 178$	Set of zones
$H = \{8, 9, \dots, 16\}$	Hourly decision grid
$i, j \in V$	Node indices
$k \in K$	Team index
$h \in H$	Decision epoch
d_{ij}	Travel time from node i to node j [min]
<i>State components</i>	
$s_h = (\tau_h, X_h, Y_h, J_h, \Pi_h)$	System state at epoch h
τ_h	Elapsed workday time [min]
X_h	Team states at epoch h
v_h^k	Current location of team k
σ_h^k	Activity status of team k
b_h^k	Remaining time budget of team k [min]
$T_{\max}$	Total working day length [min]
Y_h	Station maintenance states at epoch h
$m_h^i \in \{0, 1\}$	Periodic maintenance completion indicator
dsm_h^i	Days since last maintenance of station i
$J_h = (W_h^{\text{reg}}, W_h^{\text{fault}}, W_h^{\text{carry}})$	Open maintenance jobs
W_h^{reg}	Regular maintenance jobs
W_h^{fault}	Newly arrived fault jobs
W_h^{carry}	Carryover fault jobs from previous day
Π_h	Planned routes for both teams
<i>Actions and dynamics</i>	
$A(s_h)$	Feasible action set in state s_h
$a_h \in A(s_h)$	Chosen action at epoch h
$P(s_{h+1} \mid s_h, a_h)$	Transition function
<i>Costs and value</i>	
$R = -(C_{\text{op}} + C_{\text{dt}})$	Reward function
C_{op}	Operational cost (labor + travel)
C_{dt}	Downtime cost
n_{active}	Number of active teams on a given day

Table 6 – continued

Symbol	Description
T_{shift}	Effective shift duration [h]
km_{ij}	Distance from node i to node j [km]
p_j	Nominal charging power of station j [kW]
$\Delta\tau_j^{\text{wait}}$	Effective waiting time of fault j [h]
$c_L = 35 \text{ €/h}$	Labor cost rate
$c_F = 0.30 \text{ €/km}$	Travel cost rate
$c_{\text{dt}} = 0.50 \text{ €/kWh}$	Downtime cost rate
$Q^*(s, a)$	Optimal state-action value function
<i>Fault model</i>	
$p_{\ell, \text{base}}$	Baseline fault rate for fault type ℓ [per hour]
η^i	Station-level fault scaling factor
$f(t)$	Recovery curve
λ	Residual fault factor immediately after maintenance
τ_{rec}	Recovery period [days]

Im Rahmen der Erstellung dieser wissenschaftlichen Arbeit kamen KI-basierte Hilfsmittel

☒ zum Einsatz.

☐ nicht zum Einsatz.

Die Art und der Umfang der Nutzung sind wie folgt:

Betroffener Teil der Arbeit	KI-gestütztes Hilfsmittel	Einsatzform	Beschreibung des Einsatzes	Eigenständige Prüfung erfolgt?
Alle Kapitel (Kap. 1–9)	Perplexity (ChatGPT 5.4)	Sprachliche Präzisierung	Unterstützung bei der stilistischen Umformulierung einzelner Absätze und Passagen zur Verbesserung von Klarheit und Lesefluß. Inhaltsaussagen wurden nicht verändert.	Alle Formulierungen wurden inhaltlich geprüft und bei Bedarf angepasst.
Gliederung (Kap. 1–9)	Perplexity (ChatGPT 5.4)	Strukturierung	Generierung von Gliederungsvorschlägen als Grundlage für eigenständige Entwicklung der finalen Struktur.	Die finale Gliederung wurde eigenständig erstellt und optimiert.
Literaturrecherche (Kap. 3)	Perplexity (ChatGPT 5.4)	Generierung von Suchbegriffen	Vorschläge für Suchstrings zur Recherche in Google Scholar und wissenschaftlichen Datenbanken. Eigenständige Recherche und Bewertung aller Quellen.	Alle Quellen wurden eigenständig recherchiert, geprüft und verifiziert.
Projekt- und Experimententwicklung (Kap. 4–6)	Claude Code (Sonnet 5)	Softwareentwicklung	Unterstützung bei der Entwicklung und Optimierung der Planungs- und Simulationssoftware (Features, Bug-Fixes, Tests). Eigenständige Überprüfung durch manuelle Tests und Validierung.	Code wurde geprüft, verstanden und selbstständig ausgeführt.
Grafikerstellung (Kap. 2, 5–7)	Claude Code (Sonnet 5)	Visualisierung	Skripte für die Erstellung von passenden Grafiken zur Darstellung von Konzepten und Ergebnissen.	Grafikideen wurden eigenständig ausgedacht und überprüft.

Hiermit versichere ich, die vorliegende Arbeit selbstständig verfasst und keine anderen als die angegebenen Quellen und Hilfsmittel benutzt sowie die Zitate deutlich kenntlich gemacht zu haben.

Ich erkläre weiterhin, dass die vorliegende Arbeit in gleicher oder ähnlicher Form noch nicht im Rahmen eines anderen Prüfungsverfahrens eingereicht wurde.

Würzburg, den 22. Juli 2026

Philipp Hümmer